

System Design Solutions: Problems 26-45

Complete architectural solutions with diagrams, code examples, and implementation details.

26. Distributed Key-Value Store

Problem Overview

Design a distributed key-value data store like Redis or DynamoDB that provides high availability, horizontal scalability, and tunable consistency guarantees.

Back-of-the-Envelope Estimation

- **Operations/sec:** 1 million ops/sec (70% reads, 30% writes)
- **Data size:** 1TB total data
- **Key size:** Average 50 bytes
- **Value size:** Average 1KB
- **Total keys:** ~1 billion keys
- **Nodes:** 10 nodes initially, scalable to 100+
- **Replication factor:** 3 (for durability)
- **Storage per node:** 100GB with replication

Functional Requirements

- **FR1:** Put(key, value) - Store key-value pair
- **FR2:** Get(key) - Retrieve value by key
- **FR3:** Delete(key) - Remove key-value pair
- **FR4:** Support TTL (Time To Live) for automatic expiration
- **FR5:** Atomic operations (increment, compare-and-swap)
- **FR6:** Range queries (scan keys by prefix)

Non-Functional Requirements

- **Scalability:** Horizontal scaling to 100+ nodes
- **Availability:** 99.99% uptime (4 nines)
- **Latency:** <1ms for reads, <5ms for writes (p99)
- **Consistency:** Tunable (strong, eventual, or read-your-writes)
- **Durability:** No data loss (replicated to N nodes)
- **Partition Tolerance:** Continue operating during network partitions

High-Level Architecture

Components:


- **Client Library:** Smart client with routing logic
- **Coordinator Nodes:** Handle requests, route to correct partition
- **Storage Nodes:** Store actual data with LSM trees


- **Replication Manager:** Ensure N replicas per key
- **Membership Service:** Track node health (gossip protocol)
- **Partition Manager:** Consistent hashing for data distribution

Data Storage Choices

Component	Technology	Justification
Storage Engine	LSM Tree (RocksDB)	High write throughput, efficient compaction
Replication	Multi-master	High availability, low latency writes
Partitioning	Consistent Hashing	Even distribution, minimal data movement
Membership	Gossip Protocol	Decentralized, fault-tolerant

Architecture Diagram:





Replication Strategy (N=3):

Key "user:123" → hash → position 45°
 Primary replicas (clockwise):
 1. Node 2 (45°) ← Primary/Leader
 2. Node 3 (90°) ← Replica
 3. Node 1 (135°) ← Replica

Write Path (Quorum W=2):

1. Client → hash(key) → find replicas [N2, N3, N1]
2. Send write to all 3 replicas in parallel
3. Wait for W=2 acknowledgments
4. Return success to client
5. Async: 3rd replica eventually catches up

Read Path (Quorum R=2):

1. Client → hash(key) → find replicas [N2, N3, N1]
2. Send read to R=2 replicas
3. If values match → return
4. If values differ → read repair (return latest)

Implementation:

```

import java.security.MessageDigest;
import java.time.*;
import java.util.*;
import java.util.concurrent.*;
import java.util.stream.Collectors;

/**
 * Consistent Hashing Ring for distributed data partitioning.
 * Uses virtual nodes to ensure even distribution across physical nodes.
 */
public class ConsistentHashRing {
    private final int virtualNodesPerNode;
    private final TreeMap<Long, String> ring = new TreeMap<>();
  
```

```
private final Set<String> physicalNodes = new HashSet<>();  
  
public ConsistentHashRing(List<String> nodes, int virtualNodesPerNode)  
{  
    this.virtualNodesPerNode = virtualNodesPerNode;  
    nodes.forEach(this::addNode);  
}  
  
/**  
 * Hash a key to position on the ring (0 to 2^32-1).  
 */  
private long hash(String key) {  
    try {  
        MessageDigest md = MessageDigest.getInstance("MD5");  
        byte[] digest = md.digest(key.getBytes());  
        // Use first 4 bytes for hash value  
        return ((long)(digest[0] & 0xFF) << 24) |  
               ((long)(digest[1] & 0xFF) << 16) |  
               ((long)(digest[2] & 0xFF) << 8) |  
               ((long)(digest[3] & 0xFF));  
    } catch (Exception e) {  
        return key.hashCode() & 0xFFFFFFFFL;  
    }  
}  
  
/**  
 * Add a physical node with its virtual nodes to the ring.  
 */  
public void addNode(String node) {  
    physicalNodes.add(node);  
    for (int i = 0; i < virtualNodesPerNode; i++) {  
        String virtualKey = node + ":" + i;  
        long hashValue = hash(virtualKey);  
        ring.put(hashValue, node);  
    }  
}  
  
/**  
 * Remove a node and all its virtual nodes from the ring.  
 */  
public void removeNode(String node) {  
    physicalNodes.remove(node);  
    ring.entrySet().removeIf(entry -> entry.getValue().equals(node));  
}  
  
/**  
 * Get N responsible nodes for a key (traversing clockwise).  
 * Used for replication – returns primary + replica nodes.  
 */  
public List<String> getNodesForKey(String key, int count) {  
    if (ring.isEmpty()) {  
        return Collections.emptyList();  
    }
```

```
long hashValue = hash(key);
List<String> nodes = new ArrayList<>();
Set<String> seen = new HashSet<>();

// Get entries starting from hash position (clockwise)
SortedMap<Long, String> tailMap = ring.tailMap(hashValue);

// First check entries after hash position
for (String node : tailMap.values()) {
    if (!seen.contains(node)) {
        nodes.add(node);
        seen.add(node);
        if (nodes.size() == count) break;
    }
}

// Wrap around to beginning if needed
if (nodes.size() < count) {
    for (String node : ring.values()) {
        if (!seen.contains(node)) {
            nodes.add(node);
            seen.add(node);
            if (nodes.size() == count) break;
        }
    }
}

return nodes;
}

}

/***
 * Distributed Key-Value Store with tunable consistency.
 * Supports quorum reads/writes and automatic replication.
 */
public class DistributedKVStore {

    private final String nodeId;
    private final ConsistentHashRing hashRing;
    private final LSMTreeStorage localStorage;
    private final Map<String, NodeConnection> nodeConnections;

    private final int replicationFactor; // N
    private final int writeQuorum; // W
    private final int readQuorum; // R
    // Consistency guarantee: R + W > N ensures strong consistency

    public DistributedKVStore(String nodeId, List<String> allNodes,
                               int replicationFactor, int writeQuorum, int
readQuorum) {
        this.nodeId = nodeId;
        this.replicationFactor = replicationFactor;
        this.writeQuorum = writeQuorum;
        this.readQuorum = readQuorum;
    }
}
```

```
        this.hashRing = new ConsistentHashRing(allNodes, 150);
        this.localStorage = new LSMTreeStorage();
        this.nodeConnections = allNodes.stream()
            .collect(Collectors.toMap(n -> n, NodeConnection::new));
    }

    /**
     * Store a key-value pair with optional TTL.
     * Writes to W replicas before returning success.
     */
    public CompletableFuture<Boolean> put(String key, Object value,
Duration ttl) {
    List<String> replicas = hashRing.getNodesForKey(key,
replicationFactor);

    VersionedValue versionedValue = new VersionedValue(
        value,
        Instant.now(),
        generateVersion(),
        ttl
    );

    // Write to all replicas in parallel
    List<CompletableFuture<Boolean>> writeFutures = replicas.stream()
        .map(replica -> {
            if (replica.equals(nodeId)) {
                return localPut(key, versionedValue);
            } else {
                return remotePut(replica, key, versionedValue);
            }
        })
        .collect(Collectors.toList());

    // Wait for quorum acknowledgments
    return CompletableFuture.supplyAsync(() -> {
        int successCount = 0;
        for (CompletableFuture<Boolean> future : writeFutures) {
            try {
                if (future.get(5, TimeUnit.SECONDS)) {
                    successCount++;
                }
            } catch (Exception e) {
                // Write to this replica failed
            }
        }
        return successCount >= writeQuorum;
    });
}

/**
 * Retrieve a value by key with quorum read.
 * Performs read repair if replicas have inconsistent data.
 */
public CompletableFuture<Optional<Object>> get(String key) {
```

```
        List<String> replicas = hashRing.getNodesForKey(key,
replicationFactor);

        // Read from R replicas
        List<CompletableFuture<VersionedValue>> readFutures =
            replicas.subList(0, Math.min(readQuorum,
replicas.size())).stream()
                .map(replica -> {
                    if (replica.equals(nodeId)) {
                        return localGet(key);
                    } else {
                        return remoteGet(replica, key);
                    }
                })
                .collect(Collectors.toList());

    return CompletableFuture.supplyAsync(() -> {
        List<VersionedValue> results = new ArrayList<>();

        for (CompletableFuture<VersionedValue> future : readFutures) {
            try {
                VersionedValue value = future.get(5,
TimeUnit.SECONDS);
                if (value != null) {
                    results.add(value);
                }
            } catch (Exception e) {
                // Read from this replica failed
            }
        }

        if (results.isEmpty()) {
            return Optional.empty();
        }

        // Resolve conflicts - last-write-wins
        VersionedValue latest = results.stream()
            .max(Comparator.comparing(VersionedValue::getTimestamp))
            .orElse(null);

        // Check TTL expiration
        if (latest != null && latest.isExpired()) {
            delete(key);
            return Optional.empty();
        }

        // Read repair if versions differ
        Set<String> versions = results.stream()
            .map(VersionedValue::getVersion)
            .collect(Collectors.toSet());

        if (versions.size() > 1 && latest != null) {
            readRepair(key, latest, replicas);
        }
    });
}
```

```
        return Optional.ofNullable(latest != null ? latest.getValue()
: null);
    });
}

/***
 * Delete a key using tombstone for eventual consistency.
 */
public CompletableFuture<Boolean> delete(String key) {
    VersionedValue tombstone = new VersionedValue(
        null, Instant.now(), generateVersion(), null, true
    );
    return put(key, tombstone, null);
}

private CompletableFuture<Boolean> localPut(String key, VersionedValue
value) {
    return CompletableFuture.supplyAsync(() -> {
        try {
            localStorage.put(key, value);
            return true;
        } catch (Exception e) {
            return false;
        }
    });
}

private CompletableFuture<VersionedValue> localGet(String key) {
    return CompletableFuture.supplyAsync(() -> localStorage.get(key));
}

private CompletableFuture<Boolean> remotePut(String node, String key,
                                              VersionedValue value) {
    return nodeConnections.get(node).put(key, value);
}

private CompletableFuture<VersionedValue> remoteGet(String node,
String key) {
    return nodeConnections.get(node).get(key);
}

private void readRepair(String key, VersionedValue correctValue,
                       List<String> replicas) {
    // Asynchronously repair stale replicas
    for (String replica : replicas) {
        if (!replica.equals(nodeId)) {
            remotePut(replica, key, correctValue);
        }
    }
}

private String generateVersion() {
    return nodeId + ":" + Instant.now().toEpochMilli();
```

```
    }

}

/***
 * Versioned value with timestamp, TTL, and tombstone support.
 */
@Data
@AllArgsConstructor
public class VersionedValue {
    private Object value;
    private Instant timestamp;
    private String version;
    private Duration ttl;
    private boolean tombstone;

    public VersionedValue(Object value, Instant timestamp, String version,
Duration ttl) {
        this(value, timestamp, version, ttl, false);
    }

    public boolean isExpired() {
        if (ttl == null) return false;
        return Instant.now().isAfter(timestamp.plus(ttl));
    }
}

/***
 * LSM Tree storage engine (simplified implementation).
 * Provides high write throughput with log-structured storage.
 */
public class LSMTreeStorage {

    private final ConcurrentHashMap<String, VersionedValue> memtable =
        new ConcurrentHashMap<>();
    private final List<SSTable> sstables =
        Collections.synchronizedList(new ArrayList<>());
    private static final int MEMTABLE_SIZE_LIMIT = 1000;

    public void put(String key, VersionedValue value) {
        memtable.put(key, value);

        if (memtable.size() >= MEMTABLE_SIZE_LIMIT) {
            flushMemtable();
        }
    }

    public VersionedValue get(String key) {
        // Check memtable first (most recent data)
        VersionedValue value = memtable.get(key);
        if (value != null) return value;

        // Check SSTables from newest to oldest
        for (int i = sstables.size() - 1; i >= 0; i--) {
            value = sstables.get(i).get(key);
        }
    }
}
```

```
        if (value != null) return value;
    }

    return null;
}

private synchronized void flushMemtable() {
    if (memtable.size() < MEMTABLE_SIZE_LIMIT) return;

    // Sort entries and create SSTable
    Map<String, VersionedValue> sortedEntries = new TreeMap<>(memtable);
    SSTable sstable = new SSTable(sortedEntries);
    sstables.add(sstable);

    memtable.clear();

    // Trigger compaction if needed
    if (sstables.size() > 5) {
        compact();
    }
}

private synchronized void compact() {
    // Merge all SSTables, keeping only latest versions
    Map<String, VersionedValue> merged = new TreeMap<>();

    for (SSTable sstable : sstables) {
        for (Map.Entry<String, VersionedValue> entry : sstable.entries()) {
            String key = entry.getKey();
            VersionedValue existing = merged.get(key);
            if (existing == null ||

entry.getValue().getTimestamp().isAfter(existing.getTimestamp())) {
                merged.put(key, entry.getValue());
            }
        }
    }

    sstables.clear();
    sstables.add(new SSTable(merged));
}

/**
 * Sorted String Table - immutable on-disk storage structure.
 */
public class SSTable {
    private final Map<String, VersionedValue> data;

    public SSTable(Map<String, VersionedValue> sortedData) {
        this.data = new TreeMap<>(sortedData);
    }
}
```

```

    public VersionedValue get(String key) {
        return data.get(key);
    }

    public Set<Map.Entry<String, VersionedValue>> entries() {
        return data.entrySet();
    }
}

/**
 * Connection abstraction for remote node communication.
 */
public class NodeConnection {
    private final String nodeAddress;

    public NodeConnection(String nodeAddress) {
        this.nodeAddress = nodeAddress;
    }

    public CompletableFuture<Boolean> put(String key, VersionedValue value) {
        // RPC call implementation (gRPC, HTTP, or custom protocol)
        return CompletableFuture.completedFuture(true);
    }

    public CompletableFuture<VersionedValue> get(String key) {
        // RPC call implementation
        return CompletableFuture.completedFuture(null);
    }
}

```

Gossip Protocol for Membership:

```

import java.util.*;
import java.util.concurrent.*;
import java.time.Instant;

/**
 * Gossip protocol implementation for cluster membership management
 * and failure detection in distributed systems.
 */
public class GossipMembership {

    private final String nodeId;
    private final ConcurrentHashMap<String, NodeState> nodes = new
    ConcurrentHashMap<>();
    private final int heartbeatIntervalMs;
    private final int gossipFanout;
    private final int failureTimeoutMs;
    private final ScheduledExecutorService scheduler;
    private final Random random = new Random();

```

```
public GossipMembership(String nodeId, List<String> seedNodes,
                      int heartbeatIntervalMs, int gossipFanout,
                      int failureTimeoutMs) {
    this.nodeId = nodeId;
    this.heartbeatIntervalMs = heartbeatIntervalMs;
    this.gossipFanout = gossipFanout;
    this.failureTimeoutMs = failureTimeoutMs;
    this.scheduler = Executors.newScheduledThreadPool(2);

    // Initialize with seed nodes
    for (String seed : seedNodes) {
        nodes.put(seed, new NodeState(seed, true));
    }
    nodes.put(nodeId, new NodeState(nodeId, true));
}

/**
 * Start the gossip protocol background threads.
 */
public void start() {
    // Schedule periodic gossip
    scheduler.scheduleAtFixedRate(
        this::gossipRound,
        0,
        heartbeatIntervalMs,
        TimeUnit.MILLISECONDS
    );

    // Schedule failure detection
    scheduler.scheduleAtFixedRate(
        this::detectFailures,
        heartbeatIntervalMs,
        heartbeatIntervalMs,
        TimeUnit.MILLISECONDS
    );
}

/**
 * Execute one round of gossip protocol.
 */
private void gossipRound() {
    // Increment own heartbeat
    NodeState self = nodes.get(nodeId);
    self.incrementHeartbeat();

    // Select random nodes to gossip with
    List<String> targets = selectGossipTargets();

    // Send gossip to each target
    for (String target : targets) {
        sendGossip(target);
    }
}
```

```
/**  
 * Select random alive nodes for gossiping.  
 */  
private List<String> selectGossipTargets() {  
    List<String> aliveNodes = nodes.entrySet().stream()  
        .filter(e -> e.getValue().isAlive() &&  
!e.getKey().equals(nodeId))  
        .map(Map.Entry::getKey)  
        .collect(Collectors.toList());  
  
    Collections.shuffle(aliveNodes, random);  
  
    int count = Math.min(gossipFanout, aliveNodes.size());  
    return aliveNodes.subList(0, count);  
}  
  
/**  
 * Send gossip message to target node.  
 */  
private void sendGossip(String target) {  
    GossipMessage message = new GossipMessage(nodeId, new HashMap<>(nodes));  
  
    // Send via network (RPC call)  
    // networkClient.sendGossip(target, message);  
}  
  
/**  
 * Process received gossip message and merge state.  
 */  
public void receiveGossip(GossipMessage message) {  
    for (Map.Entry<String, NodeState> entry :  
message.getNodeStates().entrySet()) {  
        String remoteNodeId = entry.getKey();  
        NodeState remoteState = entry.getValue();  
  
        NodeState localState = nodes.get(remoteNodeId);  
  
        if (localState == null) {  
            // New node discovered  
            nodes.put(remoteNodeId, remoteState.copy());  
            System.out.println("Discovered new node: " +  
remoteNodeId);  
        } else if (remoteState.getHeartbeat() >  
localState.getHeartbeat()) {  
            // Remote has newer information  
            localState.update(remoteState);  
        }  
    }  
}  
  
/**  
 * Detect failed nodes based on heartbeat timeout.  
 */
```

```
/*
private void detectFailures() {
    long now = System.currentTimeMillis();

    for (Map.Entry<String, NodeState> entry : nodes.entrySet()) {
        String targetNodeId = entry.getKey();
        NodeState state = entry.getValue();

        if (targetNodeId.equals(nodeId)) continue;

        if (state.isAlive() &&
            (now - state.getLastUpdateMs()) > failureTimeoutMs) {
            // Node suspected as failed
            state.markAsFailed();
            System.out.println("Node " + targetNodeId + " marked as
FAILED");
            handleNodeFailure(targetNodeId);
        }
    }
}

/**
 * Handle node failure - trigger data rebalancing.
 */
private void handleNodeFailure(String failedNode) {
    // Remove from hash ring
    // Trigger data redistribution
    // Notify other components
}

/**
 * Get all alive nodes in the cluster.
 */
public List<String> getAliveNodes() {
    return nodes.entrySet().stream()
        .filter(e -> e.getValue().isAlive())
        .map(Map.Entry::getKey)
        .collect(Collectors.toList());
}

public void shutdown() {
    scheduler.shutdown();
}
}

/**
 * State of a node in the cluster.
 */
public class NodeState {
    private final String nodeId;
    private volatile boolean alive;
    private volatile long heartbeat;
    private volatile long lastUpdateMs;
```

```

public NodeState(String nodeId, boolean alive) {
    this.nodeId = nodeId;
    this.alive = alive;
    this.heartbeat = 0;
    this.lastUpdateMs = System.currentTimeMillis();
}

public synchronized void incrementHeartbeat() {
    this.heartbeat++;
    this.lastUpdateMs = System.currentTimeMillis();
}

public synchronized void update(NodeState other) {
    this.heartbeat = other.heartbeat;
    this.lastUpdateMs = System.currentTimeMillis();
    this.alive = true; // Receiving updates means node is alive
}

public synchronized void markAsFailed() {
    this.alive = false;
}

public NodeState copy() {
    NodeState copy = new NodeState(nodeId, alive);
    copy.heartbeat = this.heartbeat;
    copy.lastUpdateMs = this.lastUpdateMs;
    return copy;
}

// Getters
public String getNodeId() { return nodeId; }
public boolean isAlive() { return alive; }
public long getHeartbeat() { return heartbeat; }
public long getLastUpdateMs() { return lastUpdateMs; }
}

/**
 * Gossip message exchanged between nodes.
 */
@Data
@AllArgsConstructor
public class GossipMessage {
    private String senderId;
    private Map<String, NodeState> nodeStates;
}

```

Trade-offs & Assumptions

- **Consistency vs Availability:** Tunable quorum ($R+W>N$) allows choosing between strong consistency and high availability
- **CAP Theorem:** System is AP (Available + Partition Tolerant) by default, CP with strong quorum
- **Replication Factor:** $N=3$ balances durability and storage cost

- **Virtual Nodes:** 150 per physical node ensures even distribution
 - **Assumption:** Network partitions are rare; gossip protocol detects failures within 10 seconds
-

27. Movie Seat Booking System

Problem Overview

Design a movie ticket booking system handling concurrent seat reservations with payment integration, preventing double bookings, and managing the complete flow from seat selection to confirmation.

Back-of-the-Envelope Estimation

- **Theaters:** 500 theaters
- **Screens per theater:** 10 screens
- **Shows per screen/day:** 5 shows
- **Seats per show:** 200 seats
- **Total seats/day:** $500 \times 10 \times 5 \times 200 = 5 \text{ million seats}$
- **Booking rate:** 20% occupancy = 1 million bookings/day
- **Peak bookings/sec:** $1\text{M} / 86400 \times 10 \text{ (peak factor)} = \sim 115 \text{ bookings/sec}$
- **Concurrent users:** 10K users simultaneously booking


Functional Requirements


- **FR1:** Browse movies, theaters, and showtimes
- **FR2:** View seat map with availability in real-time
- **FR3:** Select seats and hold temporarily (10 minutes)
- **FR4:** Process payment
- **FR5:** Generate and send ticket confirmation
- **FR6:** Cancel bookings with refund
- **FR7:** Handle booking failures and rollbacks

Non-Functional Requirements

- **Scalability:** Handle 115 bookings/sec peak load
- **Availability:** 99.9% uptime
- **Latency:** <500ms for seat selection, <3s for payment
- **Consistency:** Strong consistency (no double bookings)
- **Atomicity:** Seat reservation + payment atomic

High-Level Architecture





Booking Flow with DB-Level Locking:

1. User Selects Seats

↳ Check availability (optimistic)

2. Hold Seats (Pessimistic Lock)

```
BEGIN TRANSACTION
SELECT * FROM seats
WHERE show_id = ? AND seat_numbers IN (?)
FOR UPDATE NOWAIT; -- Acquire row-level lock
```

IF all seats available:

```
UPDATE seats SET status = 'held',
held_by = session_id,
held_until = NOW() + INTERVAL '10 minutes'
WHERE ...
```

ELSE:

```
ROLLBACK -- Seats taken
```

```
COMMIT
```

3. User Enters Payment

↳ 10 min timer countdown

4. Process Payment

```
BEGIN TRANSACTION
```

-- Verify seats still held by this session

```
SELECT * FROM seats
WHERE show_id = ? AND seat_numbers IN (?)
AND held_by = session_id
FOR UPDATE;
```

```
-- Process payment
payment_result = payment_gateway.charge()

IF payment successful:
    INSERT INTO bookings (...)

    UPDATE seats SET status = 'booked'
    COMMIT

ELSE:
    UPDATE seats SET status = 'available'
    ROLLBACK
END TRANSACTION
```

5. Generate Ticket

↳ QR code, PDF, email

Double Booking Prevention Strategy:

Layer 1: Application-Level Lock

- Redis distributed lock
- Prevents concurrent hold attempts

Layer 2: Database Row Lock

- SELECT FOR UPDATE NOWAIT
- Database guarantees exclusivity

Layer 3: Unique Constraint

- UNIQUE(show_id, seat_number)
- Final safety net

Database Schema:

```
-- Movies and shows
movies (
    id BIGSERIAL PRIMARY KEY,
    title VARCHAR(255) NOT NULL,
    duration INT NOT NULL, -- minutes
    genre VARCHAR(50),
    rating VARCHAR(10)
);
```

```
theaters (
    id BIGSERIAL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    location VARCHAR(255),
    total_screens INT
);
```

```
screens (
```

```
id BIGSERIAL PRIMARY KEY,
theater_id BIGINT REFERENCES theaters(id),
screen_number INT,
total_seats INT,
seat_layout JSONB -- {"rows": 10, "cols": 20}
);

shows (
    id BIGSERIAL PRIMARY KEY,
    movie_id BIGINT REFERENCES movies(id),
    screen_id BIGINT REFERENCES screens(id),
    show_time TIMESTAMP NOT NULL,
    price_base DECIMAL(10,2),
    status VARCHAR(20) DEFAULT 'scheduled'
);

-- Seats with locking
seats (
    id BIGSERIAL PRIMARY KEY,
    show_id BIGINT REFERENCES shows(id),
    seat_row VARCHAR(5) NOT NULL,
    seat_number INT NOT NULL,
    seat_type VARCHAR(20), -- regular, premium, vip
    price DECIMAL(10,2) NOT NULL,
    status VARCHAR(20) NOT NULL DEFAULT 'available',
        -- available, held, booked, blocked
    held_by VARCHAR(100), -- session_id
    held_until TIMESTAMP,
    UNIQUE(show_id, seat_row, seat_number)
);

CREATE INDEX idx_seats_show_status ON seats(show_id, status);
CREATE INDEX idx_seats_held_until ON seats(held_until)
WHERE status = 'held';

-- Bookings
bookings (
    id UUID PRIMARY KEY,
    user_id BIGINT,
    show_id BIGINT REFERENCES shows(id),
    seat_ids BIGINT[],
    total_amount DECIMAL(10,2),
    payment_id VARCHAR(100),
    status VARCHAR(20) NOT NULL,
        -- pending, confirmed, cancelled, refunded
    booking_time TIMESTAMP DEFAULT NOW(),
    confirmation_code VARCHAR(20) UNIQUE
);

-- Payments
payments (
    id BIGSERIAL PRIMARY KEY,
    booking_id UUID REFERENCES bookings(id),
    amount DECIMAL(10,2) NOT NULL,
```

```
payment_method VARCHAR(20),
gateway_transaction_id VARCHAR(100),
status VARCHAR(20) NOT NULL,
-- initiated, success, failed, refunded
created_at TIMESTAMP DEFAULT NOW()
);
```

Booking Service Implementation:

```
import org.springframework.stereotype.Service;
import org.springframework.transaction.annotation.Transactional;
import org.springframework.data.redis.core.RedisTemplate;
import org.springframework.data.redis.core.script.DefaultRedisScript;
import java.time.*;
import java.util.*;
import java.util.concurrent.*;

/**
 * Movie ticket booking service with distributed locking and
 * database-level concurrency control for double-booking prevention.
 */
@Service
public class BookingService {

    private final JdbcTemplate db;
    private final RedisTemplate<String, String> redis;
    private final PaymentService paymentService;
    private final TicketGeneratorService ticketService;

    private static final int HOLD_DURATION_SECONDS = 600; // 10 minutes
    private static final Duration LOCK_EXPIRY = Duration.ofSeconds(30);

    // Lua script for safe lock release
    private static final String RELEASE_LOCK_SCRIPT = """
        if redis.call("GET", KEYS[1]) == ARGV[1] then
            return redis.call("DEL", KEYS[1])
        else
            return 0
        end
    """;

    @Autowired
    public BookingService(JdbcTemplate db,
                          RedisTemplate<String, String> redis,
                          PaymentService paymentService,
                          TicketGeneratorService ticketService) {
        this.db = db;
        this.redis = redis;
        this.paymentService = paymentService;
        this.ticketService = ticketService;
    }
}
```

```
/**  
 * Hold seats temporarily for a user.  
 * Uses distributed lock + database row lock for double-booking  
 prevention.  
 */  
@Transactional(isolation = Isolation.SERIALIZABLE)  
public HoldResult holdSeats(long showId, List<Long> seatIds, String  
sessionId) {  
    // Generate lock key for these specific seats  
    String lockKey = buildLockKey(showId, seatIds);  
  
    // Step 1: Acquire distributed lock (prevents concurrent attempts)  
    boolean lockAcquired = acquireDistributedLock(lockKey, sessionId);  
    if (!lockAcquired) {  
        return HoldResult.failure("CONCURRENT_BOOKING",  
            "Someone else is booking these seats");  
    }  
  
    try {  
        // Step 2: Database transaction with row-level locking  
        // NOWAIT fails immediately if rows are locked  
        String selectQuery = """";  
        SELECT id, status, held_by, held_until  
        FROM seats  
        WHERE show_id = ? AND id = ANY(?)  
        FOR UPDATE NOWAIT  
        """;  
  
        List<SeatRow> seats;  
        try {  
            seats = db.query(selectQuery,  
                new SeatRowMapper(), showId, seatIds.toArray(new  
Long[0]));  
        } catch (CannotAcquireLockException e) {  
            return HoldResult.failure("SEATS_LOCKED",  
                "Seats are being booked by another user");  
        }  
  
        // Step 3: Validate all seats are available  
        for (SeatRow seat : seats) {  
            if ("booked".equals(seat.getStatus())) {  
                return HoldResult.failure("SEATS_ALREADY_BOOKED",  
                    "Seat " + seat.getId() + " is already booked");  
            }  
  
            if ("held".equals(seat.getStatus()) &&  
                !sessionId.equals(seat.getHeldBy())) {  
                // Check if hold has expired  
                if (seat.getHeldUntil() != null &&  
                    seat.getHeldUntil().isAfter(Instant.now())) {  
                    return HoldResult.failure("SEATS_HELD",  
                        "Seat " + seat.getId() + " is held by another  
user");  
                }  
            }  
        }  
    }  
}
```

```
        }

    }

    // Step 4: Hold the seats
    Instant heldUntil =
Instant.now().plusSeconds(HOLD_DURATION_SECONDS);

    String updateQuery = """
        UPDATE seats
        SET status = 'held',
            held_by = ?,
            held_until = ?
        WHERE show_id = ? AND id = ANY(?)
        """;

    db.update(updateQuery, sessionId,
        Timestamp.from(heldUntil), showId, seatIds.toArray(new
Long[0]));

    return HoldResult.success(heldUntil, HOLD_DURATION_SECONDS);
}

} finally {
    // Always release distributed lock
    releaseDistributedLock(lockKey, sessionId);
}
}

/**
 * Confirm booking with payment processing.
 * Implements Saga pattern for distributed transaction handling.
 */
@Transactional(isolation = Isolation.SERIALIZABLE)
public BookingResult confirmBooking(long showId, List<Long> seatIds,
        String sessionId, long userId,
        PaymentDetails paymentDetails) {
    String bookingId = UUID.randomUUID().toString();

    try {
        // Step 1: Verify seats still held by this session
        String verifyQuery = """
            SELECT id, price, status, held_by
            FROM seats
            WHERE show_id = ? AND id = ANY(?)
            FOR UPDATE
            """;

        List<SeatRow> seats = db.query(verifyQuery,
            new SeatRowMapper(), showId, seatIds.toArray(new
Long[0]));

        // Validate ownership and status
        for (SeatRow seat : seats) {
            if (!"held".equals(seat.getStatus()) ||
                !sessionId.equals(seat.getHeldBy())) {
```

```
        throw new BookingException("SEATS_NOT_HELD",
            "Seats no longer held by this session");
    }

    BigDecimal totalAmount = seats.stream()
        .map(SeatRow::getPrice)
        .reduce(BigDecimal.ZERO, BigDecimal::add);

    // Step 2: Create pending booking record
    String insertBookingQuery = """
        INSERT INTO bookings
        (id, user_id, show_id, seat_ids, total_amount, status)
        VALUES (?, ?, ?, ?, ?, ?, 'pending')
        RETURNING confirmation_code
    """;

    String confirmationCode =
    db.queryForObject(insertBookingQuery,
        String.class, bookingId, userId, showId,
        seatIds.toArray(new Long[0]), totalAmount);

    // Step 3: Process payment (external service call)
    PaymentResult paymentResult;
    try {
        paymentResult = paymentService.charge(
            totalAmount, "USD", paymentDetails,
            Map.of("booking_id", bookingId));

        if (!paymentResult.isSuccess()) {
            throw new PaymentException(paymentResult.getError());
        }
    } catch (PaymentException e) {
        // Payment failed - rollback
        rollbackBooking(bookingId, showId, seatIds);
        return BookingResult.failure("PAYMENT_FAILED",
e.getMessage());
    }

    // Step 4: Finalize booking
    // Update seats to booked
    String updateSeatsQuery = """
        UPDATE seats
        SET status = 'booked',
            held_by = NULL,
            held_until = NULL
        WHERE show_id = ? AND id = ANY(?)
    """;
    db.update(updateSeatsQuery, showId, seatIds.toArray(new
Long[0]));

    // Update booking status
    String updateBookingQuery = """
        UPDATE bookings
    """;
```

```
        SET status = 'confirmed',
            payment_id = ?
        WHERE id = ?
        """;
    db.update(updateBookingQuery,
paymentResult.getTransactionId(), bookingId);

        // Record payment
        String insertPaymentQuery = """
            INSERT INTO payments
            (booking_id, amount, payment_method,
gateway_transaction_id, status)
            VALUES (?, ?, ?, ?, 'success')
            """;
    db.update(insertPaymentQuery, bookingId, totalAmount,
        paymentDetails.getMethod(),
paymentResult.getTransactionId());

        // Step 5: Generate ticket asynchronously
        CompletableFuture.runAsync(() ->
            ticketService.generateAndSendTicket(bookingId, userId));
    }

    return BookingResult.success(bookingId, confirmationCode,
totalAmount);

} catch (BookingException e) {
    return BookingResult.failure(e.getCode(), e.getMessage());
} catch (Exception e) {
    rollbackBooking(bookingId, showId, seatIds);
    throw e;
}
}

/**
 * Rollback booking in case of failure.
 */
private void rollbackBooking(String bookingId, long showId, List<Long>
seatIds) {
    // Release seats
    String releaseSeatsQuery = """
        UPDATE seats
        SET status = 'available',
            held_by = NULL,
            held_until = NULL
        WHERE show_id = ? AND id = ANY(?)
        """;
    db.update(releaseSeatsQuery, showId, seatIds.toArray(new
Long[0]));
}

// Mark booking as cancelled
String cancelBookingQuery = """
    UPDATE bookings
    SET status = 'cancelled'
    WHERE id = ?
```

```
        """";
        db.update(cancelBookingQuery, bookingId);
    }

    /**
     * Background job to release expired seat holds.
     * Should be scheduled to run every minute.
     */
    @Scheduled(fixedRate = 60000)
    @Transactional
    public void releaseExpiredHolds() {
        String query = """
            UPDATE seats
            SET status = 'available',
                held_by = NULL,
                held_until = NULL
            WHERE status = 'held' AND held_until < NOW()
            RETURNING show_id, id
        """;

        List<Map<String, Object>> released = db.queryForList(query);

        if (!released.isEmpty()) {
            System.out.println("Released " + released.size() + " expired
seat holds");
        }
    }

    private String buildLockKey(long showId, List<Long> seatIds) {
        List<Long> sorted = new ArrayList<>(seatIds);
        Collections.sort(sorted);
        return String.format("booking:show:%d:seats:%s",
            showId,
            sorted.stream().map(String::valueOf).collect(Collectors.joining(",")));
    }

    private boolean acquireDistributedLock(String key, String value) {
        Boolean result = redis.opsForValue()
            .setIfAbsent(key, value, LOCK_EXPIRY);
        return Boolean.TRUE.equals(result);
    }

    private void releaseDistributedLock(String key, String value) {
        DefaultRedisScript<Long> script = new DefaultRedisScript<>(
            RELEASE_LOCK_SCRIPT, Long.class);
        redis.execute(script, Collections.singletonList(key), value);
    }
}

/**
 * Result of seat hold operation.
 */
@Data
@Builder
```

```
public class HoldResult {  
    private boolean success;  
    private String errorCode;  
    private String message;  
    private Instant heldUntil;  
    private int holdDurationSeconds;  
  
    public static HoldResult success(Instant heldUntil, int duration) {  
        return HoldResult.builder()  
            .success(true)  
            .heldUntil(heldUntil)  
            .holdDurationSeconds(duration)  
            .build();  
    }  
  
    public static HoldResult failure(String code, String message) {  
        return HoldResult.builder()  
            .success(false)  
            .errorCode(code)  
            .message(message)  
            .build();  
    }  
}  
  
/**  
 * Result of booking confirmation.  
 */  
@Data  
@Builder  
public class BookingResult {  
    private boolean success;  
    private String errorCode;  
    private String message;  
    private String bookingId;  
    private String confirmationCode;  
    private BigDecimal totalAmount;  
  
    public static BookingResult success(String bookingId, String confirmationCode,  
                                       BigDecimal totalAmount) {  
        return BookingResult.builder()  
            .success(true)  
            .bookingId(bookingId)  
            .confirmationCode(confirmationCode)  
            .totalAmount(totalAmount)  
            .build();  
    }  
  
    public static BookingResult failure(String code, String message) {  
        return BookingResult.builder()  
            .success(false)  
            .errorCode(code)  
            .message(message)  
            .build();  
    }  
}
```

```
}
```

Trade-offs & Assumptions

- **Pessimistic Locking:** `SELECT FOR UPDATE NOWAIT` prevents concurrent bookings but reduces concurrency
 - **Hold Duration:** 10 minutes balances user experience vs inventory blocking
 - **Saga Pattern:** Compensating transactions handle payment failures
 - **Assumption:** 80% of holds result in successful bookings
 - **Distributed Lock:** Redis lock prevents concurrent hold attempts before DB lock
-

28. E-commerce Top Sellers

Problem Overview

Design a system to track and display top 20-30 selling items in real-time for an e-commerce platform with millions of products and high transaction volume.

Back-of-the-Envelope Estimation

- **Products:** 10 million SKUs
- **Orders/day:** 5 million orders
- **Items per order:** Average 2.5 items
- **Total item purchases/day:** 12.5 million
- **Peak purchases/sec:** $12.5M / 86400 \times 10 = \sim 1,450/\text{sec}$
- **Top sellers update frequency:** Every 5 minutes
- **Categories:** 1,000 categories


Functional Requirements

- **FR1:** Track item purchase counts in real-time
- **FR2:** Calculate top 20-30 sellers globally
- **FR3:** Calculate top sellers per category
- **FR4:** Update rankings every 5 minutes
- **FR5:** Display trending items (velocity-based)
- **FR6:** Historical trending data (daily, weekly, monthly)

Non-Functional Requirements

- **Scalability:** Handle 1,450 purchases/sec
- **Latency:** <100ms for fetching top sellers
- **Accuracy:** 99% accuracy acceptable (approximate counts)
- **Freshness:** Rankings updated every 5 minutes

Architecture

Purchase Event Flow:

```

ClickHouse
(Historical Data)
- Daily aggregates
- Weekly aggregates
- Trending analysis

```

Redis Data Structures:

```

ZADD top:global:current product:123 5420
ZADD top:category:electronics product:111 2103
ZADD trending:velocity product:456 892 # purchases/hour

```

Query API:

```

GET /api/top-sellers?limit=30
→ ZREVRANGE top:global:current 0 29 WITHSCORES

```

```

GET /api/top-sellers/category/electronics?limit=20
→ ZREVRANGE top:category:electronics 0 19 WITHSCORES

```

Stream Processing (Flink):

```

import org.apache.flink.api.common.functions.MapFunction;
import org.apache.flink.api.java.tuple.Tuple3;
import org.apache.flink.streaming.api.datastream.DataStream;
import
org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
import org.apache.flink.streaming.api.windowing.time.Time;
import org.apache.flink.streaming.api.windowing.windows.TimeWindow;
import org.apache.flink.streaming.connectors.kafka.FlinkKafkaConsumer;
import redis.clients.jedis.Jedis;

/**
 * Flink streaming job for real-time top sellers calculation.
 * Processes purchase events and updates Redis sorted sets.
 */
public class TopSellersStreamJob {

    public static void main(String[] args) throws Exception {
        StreamExecutionEnvironment env =
            StreamExecutionEnvironment.getExecutionEnvironment();
        env.setParallelism(3);

        // Source: Kafka purchase events
        Properties kafkaProps = new Properties();
        kafkaProps.setProperty("bootstrap.servers", "kafka:9092");
        kafkaProps.setProperty("group.id", "top-sellers-consumer");

        FlinkKafkaConsumer<PurchaseEvent> kafkaSource = new
        FlinkKafkaConsumer<>(
            "purchase-events",
            new PurchaseEventDeserializer(),

```

```
        kafkaProps
    );

    DataStream<PurchaseEvent> purchases = env.addSource(kafkaSource);

    // Extract product info and count purchases
    DataStream<Tuple3<String, String, Integer>> productCounts =
purchases
        .map(new MapFunction<PurchaseEvent, Tuple3<String, String,
Integer>>() {
            @Override
            public Tuple3<String, String, Integer> map(PurchaseEvent
event) {
                return Tuple3.of(event.getProductId(),
event.getCategory(), 1);
            }
        })
        .keyBy(tuple -> tuple.f0) // Key by product_id
        .timeWindow(Time.minutes(5)) // 5-minute tumbling window
        .sum(2); // Sum the counts

    // Update Redis sorted sets
    productCounts.addSink(new RedisTopSellersSink());

    env.execute("Top Sellers Stream Processing");
}
}

/***
 * Flink sink that updates Redis sorted sets for top sellers.
 */
public class RedisTopSellersSink extends RichSinkFunction<Tuple3<String,
String, Integer>> {

    private transient Jedis redis;

    @Override
    public void open(Configuration parameters) {
        redis = new Jedis("redis", 6379);
    }

    @Override
    public void invoke(Tuple3<String, String, Integer> value, Context
context) {
        String productId = value.f0;
        String category = value.f1;
        int count = value.f2;

        // Update global top sellers
        redis.zincrby("top:global:current", count, "product:" +
productId);

        // Update category top sellers
        redis.zincrby("top:category:" + category, count, "product:" +
```

```

productId);

        // Calculate and store velocity (trending score)
        double velocity = calculateVelocity(productId);
        redis.zadd("trending:velocity", velocity, "product:" + productId);
    }

    private double calculateVelocity(String productId) {
        // Get current hour's count
        String currentKey = "count:hour:current:" + productId;
        String previousKey = "count:hour:previous:" + productId;

        String currentStr = redis.get(currentKey);
        String previousStr = redis.get(previousKey);

        double currentCount = currentStr != null ?
Double.parseDouble(currentStr) : 0;
        double previousCount = previousStr != null ?
Double.parseDouble(previousStr) : 1;

        // Velocity = current / previous (>1 means trending up)
        return currentCount / previousCount;
    }

    @Override
    public void close() {
        if (redis != null) {
            redis.close();
        }
    }
}

/***
 * Purchase event data class.
 */
@Data
@AllArgsConstructor
@NoArgsConstructor
public class PurchaseEvent {
    private String productId;
    private String category;
    private String userId;
    private BigDecimal amount;
    private Instant timestamp;
}

```

API Service:

```

import org.springframework.web.bind.annotation.*;
import org.springframework.data.redis.core.RedisTemplate;
import org.springframework.data.redis.core.ZSetOperations;
import java.util.*;

```

```
import java.util.stream.Collectors;

/**
 * REST API for retrieving top sellers and trending products.
 */
@RestController
@RequestMapping("/api")
public class TopSellersController {

    private final TopSellersService topSellersService;

    @Autowired
    public TopSellersController(TopSellersService topSellersService) {
        this.topSellersService = topSellersService;
    }

    @GetMapping("/top-sellers")
    public List<TopSellerDto> getTopSellers(
        @RequestParam(defaultValue = "30") int limit,
        @RequestParam(required = false) String category) {
        return topSellersService.getTopSellers(limit, category);
    }

    @GetMapping("/trending")
    public List<TrendingProductDto> getTrending(
        @RequestParam(defaultValue = "20") int limit) {
        return topSellersService.getTrendingProducts(limit);
    }
}

/**
 * Service for retrieving top sellers from Redis.
 */
@Service
public class TopSellersService {

    private final RedisTemplate<String, String> redisTemplate;
    private final ProductRepository productRepository;

    @Autowired
    public TopSellersService(RedisTemplate<String, String> redisTemplate,
                           ProductRepository productRepository) {
        this.redisTemplate = redisTemplate;
        this.productRepository = productRepository;
    }

    /**
     * Get top selling products globally or by category.
     */
    public List<TopSellerDto> getTopSellers(int limit, String category) {
        // Determine Redis key
        String key = (category != null)
            ? "top:category:" + category
            : "top:global:current";
        ...
    }
}
```

```
// Get from Redis sorted set (highest scores first)
ZSetOperations<String, String> zSetOps =
redisTemplate.opsForZSet();
Set<ZSetOperations.TypedTuple<String>> results =
zSetOps.reverseRangeWithScores(key, 0, limit - 1);

if (results == null || results.isEmpty()) {
    return Collections.emptyList();
}

// Extract product IDs
List<Long> productIds = results.stream()
    .map(tuple -> {
        String productKey = tuple.getValue();
        return Long.parseLong(productKey.split(":")[1]);
    })
    .collect(Collectors.toList());

// Fetch product details from database
Map<Long, Product> productMap = productRepository
    .findAllById(productIds).stream()
    .collect(Collectors.toMap(Product::getId, p -> p));

// Build response with ranking
List<TopSellerDto> topSellers = new ArrayList<>();
int rank = 1;

for (ZSetOperations.TypedTuple<String> tuple : results) {
    String productKey = tuple.getValue();
    Long productId = Long.parseLong(productKey.split(":")[1]);
    Double score = tuple.getScore();

    Product product = productMap.get(productId);
    if (product != null) {
        topSellers.add(TopSellerDto.builder()
            .productId(productId)
            .name(product.getName())
            .price(product.getPrice())
            .imageUrl(product.getImageUrl())
            .category(product.getCategory())
            .purchaseCount(score.intValue())
            .rank(rank++)
            .build());
    }
}

return topSellers;
}

/**
 * Get trending products based on velocity (purchase acceleration).
 */
public List<TrendingProductDto> getTrendingProducts(int limit) {
```

```
        ZSetOperations<String, String> zSetOps =
redisTemplate.opsForZSet();
        Set<ZSetOperations.TypedTuple<String>> results =
zSetOps.reverseRangeWithScores("trending:velocity", 0, limit -
1);

        if (results == null || results.isEmpty()) {
            return Collections.emptyList();
        }

        List<Long> productIds = results.stream()
            .map(tuple -> Long.parseLong(tuple.getValue().split(":")[1]))
            .collect(Collectors.toList());

        Map<Long, Product> productMap = productRepository
            .findAllById(productIds).stream()
            .collect(Collectors.toMap(Product::getId, p -> p));

        List<TrendingProductDto> trending = new ArrayList<>();

        for (ZSetOperations.TypedTuple<String> tuple : results) {
            Long productId = Long.parseLong(tuple.getValue().split(":") [1]);
            Double velocity = tuple.getScore();

            Product product = productMap.get(productId);
            if (product != null) {
                trending.add(TrendingProductDto.builder()
                    .productId(productId)
                    .name(product.getName())
                    .price(product.getPrice())
                    .imageUrl(product.getImageUrl())
                    .velocity(velocity)
                    .trend(velocity > 2.0 ? "hot" : "rising")
                    .build());
            }
        }

        return trending;
    }
}

/**
 * DTO for top seller product.
 */
@Data
@Builder
public class TopSellerDto {
    private Long productId;
    private String name;
    private BigDecimal price;
    private String imageUrl;
    private String category;
    private int purchaseCount;
```

```

        private int rank;
    }

    /**
     * DTO for trending product.
     */
    @Data
    @Builder
    public class TrendingProductDto {
        private Long productId;
        private String name;
        private BigDecimal price;
        private String imageUrl;
        private double velocity;
        private String trend; // "hot" or "rising"
    }
}

```

Trade-offs & Assumptions

- **Approximate Counts:** Redis sorted sets provide fast queries but approximate counts acceptable
 - **Update Frequency:** 5-minute windows balance freshness vs computational cost
 - **Top K Size:** Tracking top 30 globally, top 20 per category limits memory usage
 - **Trending Algorithm:** Velocity-based (current/previous) simple but effective
 - **Assumption:** 80% of sales come from 20% of products (Pareto principle)
-

[Continuing with remaining problems 29-45...]

29. Multi-Datacenter Replication

Problem Overview

Design a strategy for replicating data across multiple datacenters globally while maintaining consistency, minimizing latency, and handling network partitions.

Replication Strategies

1. Master-Slave Replication:

```

Primary DC (US-East)
  ↓ Async Replication
Secondary DCs (EU, APAC)

```

Writes: US-East only
 Reads: Any DC (stale reads acceptable)
 Failover: Promote secondary to primary

2. Multi-Master Replication:

```
US-East ↔ EU ↔ APAC
    ↓      ↓      ↓
Bidirectional replication
All DCs accept writes
Conflict resolution required
```

3. Consensus-Based (Raft/Paxos):

```
Quorum writes across DCs
Majority must acknowledge
Strong consistency guarantee
Higher latency for writes
```

Implementation Considerations:

- Conflict resolution (Last-Write-Wins, CRDTs)
- Cross-DC latency (50-200ms)
- Network partition handling
- Bandwidth optimization (delta sync)

30. SIM Card Store System

Problem Overview

Design a system for a SIM card store to generate unpredictable phone numbers, track sold/unused inventory, and enable efficient retrieval.

Key Requirements

- Generate random phone numbers (no patterns)
- Track status: available, sold, reserved
- Efficient search and allocation
- Prevent number prediction

Schema:

```
phone_numbers (
    id BIGSERIAL PRIMARY KEY,
    phone_number VARCHAR(15) UNIQUE NOT NULL,
    status VARCHAR(20) NOT NULL,
    allocated_at TIMESTAMP,
    customer_id BIGINT,
    random_seed BYTEA -- for generation
);
```

```
CREATE INDEX idx_phone_status ON phone_numbers(status)
WHERE status = 'available';
```

Number Generation:

```
import java.security.SecureRandom;
import java.security.MessageDigest;
import java.math.BigInteger;

/**
 * Cryptographically secure phone number generator.
 * Produces unpredictable numbers using SecureRandom and SHA-256.
 */
public class PhoneNumberGenerator {

    private final SecureRandom secureRandom;

    public PhoneNumberGenerator() {
        this.secureRandom = new SecureRandom();
    }

    /**
     * Generate a cryptographically random phone number.
     * Uses SecureRandom for unpredictability.
     */
    public String generatePhoneNumber() {
        try {
            // Generate 8 random bytes
            byte[] randomBytes = new byte[8];
            secureRandom.nextBytes(randomBytes);

            // Hash for additional randomness and to prevent patterns
            MessageDigest digest = MessageDigest.getInstance("SHA-256");
            byte[] hash = digest.digest(randomBytes);

            // Convert to a large number
            BigInteger hashValue = new BigInteger(1, hash);

            // Mod to get 10-digit number
            BigInteger tenDigits = BigInteger.valueOf(10_000_000_000L);
            long phoneDigits = hashValue.mod(tenDigits).longValue();

            // Format with country code
            return String.format("+1%010d", phoneDigits);

        } catch (Exception e) {
            throw new RuntimeException("Phone number generation failed",
e);
        }
    }

    /**

```

```

    * Generate a batch of unique phone numbers.
    */
public List<String> generateBatch(int count) {
    Set<String> numbers = new HashSet<>();
    while (numbers.size() < count) {
        numbers.add(generatePhoneNumber());
    }
    return new ArrayList<>(numbers);
}
}

```

31-32. Optimizing Hotel Search Results & Ranking Algorithm

Search Optimization Strategies

1. Inverted Index:

```

hotels_by_location[
    "New York": [hotel1, hotel2, ...],
    "San Francisco": [hotel3, hotel4, ...]
]

hotels_by_amenity[
    "pool": [hotel1, hotel5, ...],
    "wifi": [hotel2, hotel3, ...]
]

```

2. Ranking Algorithm:

```

import java.util.*;

/**
 * Hotel ranking algorithm considering multiple factors:
 * price, rating, distance, and amenity match.
 */
public class HotelRankingService {

    // Weight factors for scoring components
    private static final double PRICE_WEIGHT = 0.3;
    private static final double RATING_WEIGHT = 0.4;
    private static final double DISTANCE_WEIGHT = 0.2;
    private static final double AMENITY_WEIGHT = 0.1;

    /**
     * Calculate composite score for a hotel based on user preferences.
     */
    public double calculateHotelScore(Hotel hotel, UserPreferences
userPrefs) {

```

```
    double score = 0.0;

    // Price score (inverse - cheaper is better, normalized)
    double priceScore = 1.0 / (hotel.getPrice() / 100.0 + 1.0);
    score += priceScore * PRICE_WEIGHT;

    // Rating score (normalized to 0-1)
    double ratingScore = hotel.getRating() / 5.0;
    score += ratingScore * RATING_WEIGHT;

    // Distance score (closer is better)
    double distanceScore = 1.0 / (hotel.getDistanceKm() + 1.0);
    score += distanceScore * DISTANCE_WEIGHT;

    // Amenity match score
    double amenityScore = calculateAmenityMatch(
        hotel.getAmenities(), userPrefs.getDesiredAmenities());
    score += amenityScore * AMENITY_WEIGHT;

    return score;
}

/**
 * Calculate amenity match percentage.
 */
private double calculateAmenityMatch(Set<String> hotelAmenities,
                                      Set<String> desiredAmenities) {
    if (desiredAmenities == null || desiredAmenities.isEmpty()) {
        return 1.0; // No preferences means full match
    }

    Set<String> intersection = new HashSet<>(hotelAmenities);
    intersection.retainAll(desiredAmenities);

    return (double) intersection.size() / desiredAmenities.size();
}

/**
 * Rank hotels by composite score.
 */
public List<Hotel> rankHotels(List<Hotel> hotels, UserPreferences
userPrefs) {
    return hotels.stream()
        .peek(hotel -> hotel.setScore(calculateHotelScore(hotel,
userPrefs)))

        .sorted(Comparator.comparingDouble(Hotel::getScore).reversed())
            .collect(Collectors.toList());
}

/**
 * Hotel entity with scoring support.
*/
```

```

@Data
public class Hotel {
    private Long id;
    private String name;
    private double price;
    private double rating;
    private double distanceKm;
    private Set<String> amenities;
    private double score; // Calculated score for ranking
}

/**
 * User search preferences.
 */
@Data
public class UserPreferences {
    private Double maxPrice;
    private Integer minRating;
    private Double maxDistanceKm;
    private Set<String> desiredAmenities;
    private String priceRange; // "budget", "mid", "luxury"
}

```

3. Elasticsearch Query:

```
{
  "query": {
    "function_score": {
      "query": {
        "bool": {
          "must": [
            {"geo_distance": {
              "distance": "10km",
              "location": {"lat": 40.7, "lon": -74.0}
            }},
            {"range": {"price": {"lte": 200}}}
          ]
        }
      },
      "functions": [
        {"field_value_factor": {
          "field": "rating",
          "factor": 2.0
        }},
        {"gauss": {
          "price": {
            "origin": 100,
            "scale": 50
          }
        }}
      ]
    }
  }
}
```


```
}
```

33. Real-time Chat Application

Problem Overview

Design a real-time chat application supporting one-to-one and group messaging with typing indicators, read receipts, and message history.

Architecture



Message Schema:

```
messages (
    conversation_id UUID,
    message_id TIMEUUID,
    sender_id UUID,
    content TEXT,
    timestamp TIMESTAMP,
    status VARCHAR(20), -- sent, delivered, read
    PRIMARY KEY (conversation_id, message_id)
) WITH CLUSTERING ORDER BY (message_id DESC);
```

WebSocket Handler:

```

io.on('connection', (socket) => {
  socket.on('send_message', async (data) => {
    const message = {
      id: uuidv1(),
      conversation_id: data.conversation_id,
      sender_id: socket.user_id,
      content: data.content,
      timestamp: new Date()
    };

    // Store in Cassandra
    await cassandra.insert('messages', message);

    // Publish via Redis Pub/Sub
    redis.publish(`chat:${data.conversation_id}`,
      JSON.stringify(message));

    // Emit to recipients
    io.to(data.conversation_id).emit('new_message', message);
  });
});

```

34. Distributed Message Broker

Problem Overview

Design a high-throughput distributed message broker like Kafka for pub/sub messaging with guaranteed delivery and ordering.

Architecture

Producers → Brokers (Partitioned Topics) → Consumers

Topic: orders

- └─ Partition 0 → Consumer Group A
- └─ Partition 1 → Consumer Group A
- └─ Partition 2 → Consumer Group A

Features:

- Partitioning for parallelism
- Replication for durability
- Consumer groups for load balancing
- Offset management for exactly-once

Key Design Decisions:

- Log-based storage (append-only)
- Partition assignment (consistent hashing)

- Replication (ISR - In-Sync Replicas)
 - Consumer offset commits
-

35. Cloud File Storage (Dropbox)

Problem Overview

Design a cloud file storage service with sync, sharing, and version control (covered in detail in Problem #9).

Key Features:

- Chunking (4MB blocks)
 - Deduplication (block-level)
 - Delta sync (only changed chunks)
 - Conflict resolution (version vectors)
-

36. Distributed Configuration Store

Problem Overview

Design a distributed configuration store like etcd or ZooKeeper for storing application config with strong consistency.

Features

- Key-value storage
- Watch mechanism (notifications on changes)
- TTL support
- Transactions
- Consensus (Raft algorithm)

Use Cases:

- Service discovery
- Leader election
- Distributed locks
- Feature flags

```
import java.util.concurrent.*;
import java.util.function.Consumer;

/**
 * Distributed configuration store client interface.
 * Provides key-value storage with watch capability for configuration
management.
 */
public interface DistributedConfigClient {
```

```
* Set a configuration value.  
*/  
CompletableFuture<Void> set(String key, String value);  
  
/**  
 * Get a configuration value.  
 */  
CompletableFuture<String> get(String key);  
  
/**  
 * Delete a configuration key.  
 */  
CompletableFuture<Void> delete(String key);  
  
/**  
 * Watch for changes on a key prefix.  
 * Callback is invoked whenever a key matching the prefix is modified.  
 */  
void watch(String keyPrefix, Consumer<ConfigChangeEvent> callback);  
}  
  
/**  
 * Configuration change event.  
 */  
@Data  
@AllArgsConstructor  
public class ConfigChangeEvent {  
    private String key;  
    private String value;  
    private ChangeType type; // PUT, DELETE  
  
    public enum ChangeType {  
        PUT, DELETE  
    }  
}  
  
// Example usage  
public class ConfigurationManager {  
  
    private final DistributedConfigClient configStore;  
  
    public ConfigurationManager(DistributedConfigClient configStore) {  
        this.configStore = configStore;  
    }  
  
    public void setupConfiguration() {  
        // Set configuration value  
        configStore.set("/app/db_host", "db.example.com")  
            .thenRun(() -> System.out.println("Configuration saved"));  
  
        // Watch for changes on /app/ prefix  
        configStore.watch("/app/", event -> {  
            System.out.printf("Config changed: %s = %s (type: %s)%n",  
                event.getKey(), event.getValue(), event.getType());  
        });  
    }  
}
```

```
// React to configuration changes
if (event.getKey().equals("/app/db_host")) {
    reconnectDatabase(event.getValue());
}
});

private void reconnectDatabase(String newHost) {
    // Handle database reconnection
}
}
```

37. Nearby Places Recommender (Yelp)

Problem Overview

Design a location-based service to find nearby restaurants/businesses with ratings and reviews (similar to Problem #6 - Proximity Search).

Geospatial Indexing:

- Quadtree partitioning
- Geohash (precision-based cells)
- PostGIS (R-tree index)
- Elasticsearch geo queries

Ranking:

```
import java.util.*;
import java.util.stream.Collectors;

/**
 * Service for ranking nearby places (restaurants, businesses)
 * based on distance, rating, popularity, and user preferences.
 */
public class PlaceRankingService {

    /**
     * Rank places by weighted score based on multiple factors.
     */
    public List<Place> rankPlaces(List<Place> places,
                                    GeoLocation userLocation,
                                    UserPreferences userPrefs) {

        return places.stream()
            .peek(place -> {
                double score = 0.0;
```

```
// Distance penalty (40% weight) - closer is better
    double distance = calculateHaversineDistance(
        userLocation, place.getLocation());
    double distanceScore = 1.0 / (distance + 0.1);
    score += distanceScore * 0.4;

    // Rating (30% weight)
    double ratingScore = place.getRating() / 5.0;
    score += ratingScore * 0.3;

    // Review count / popularity (20% weight)
    double reviewScore = Math.min(place.getReviewCount() /
1000.0, 1.0);
    score += reviewScore * 0.2;

    // Price range match (10% weight)
    if (userPrefs != null &&
place.getPriceRange().equals(userPrefs.getPreferredPriceRange())) {
        score += 0.1;
    }

    place.setScore(score);
}

.sorted(Comparator.comparingDouble(Place::getScore).reversed())
.collect(Collectors.toList());
}

/**
 * Calculate distance between two points using Haversine formula.
 */
private double calculateHaversineDistance(GeoLocation loc1,
GeoLocation loc2) {
    final double R = 6371.0; // Earth's radius in km

    double lat1Rad = Math.toRadians(loc1.getLatitude());
    double lat2Rad = Math.toRadians(loc2.getLatitude());
    double deltaLat = Math.toRadians(loc2.getLatitude() -
loc1.getLatitude());
    double deltaLon = Math.toRadians(loc2.getLongitude() -
loc1.getLongitude());

    double a = Math.sin(deltaLat / 2) * Math.sin(deltaLat / 2) +
Math.cos(lat1Rad) * Math.cos(lat2Rad) *
Math.sin(deltaLon / 2) * Math.sin(deltaLon / 2);

    double c = 2 * Math.atan2(Math.sqrt(a), Math.sqrt(1 - a));

    return R * c;
}
```

```

/**
 * Place entity (restaurant, business, etc.).
 */
@Data
public class Place {
    private Long id;
    private String name;
    private GeoLocation location;
    private double rating;
    private int reviewCount;
    private String priceRange; // "$", "$$", "$$$", "$$$$"
    private List<String> categories;
    private double score; // Computed ranking score
}

@Data
@AllArgsConstructor
public class GeoLocation {
    private double latitude;
    private double longitude;
}

```

38. Gaming Leaderboard

Problem Overview

Design a real-time gaming leaderboard supporting millions of players with efficient rank queries and updates.

Architecture

```

Player Score Update:
ZADD leaderboard:global player_id score

Get Player Rank:
ZREVRANK leaderboard:global player_id

Get Top 100:
ZREVRANGE leaderboard:global 0 99 WITHSCORES

Get Nearby Players (player's rank ±10):
rank = ZREVRANK leaderboard:global player_id
ZREVRANGE leaderboard:global (rank-10) (rank+10) WITHSCORES

```

Redis Sorted Set Operations:

```

import org.springframework.data.redis.core.RedisTemplate;
import org.springframework.data.redis.core.ZSetOperations;

```

```
import org.springframework.stereotype.Service;
import java.util.*;
import java.util.stream.Collectors;

/**
 * Real-time gaming leaderboard using Redis Sorted Sets.
 * Supports efficient score updates and rank queries for millions of
players.
 */
@Service
public class Leaderboard {

    private final RedisTemplate<String, String> redisTemplate;
    private final ZSetOperations<String, String> zSetOps;
    private static final String LEADERBOARD_KEY = "leaderboard:global";

    @Autowired
    public Leaderboard(RedisTemplate<String, String> redisTemplate) {
        this.redisTemplate = redisTemplate;
        this.zSetOps = redisTemplate.opsForZSet();
    }

    /**
     * Update a player's score.
     * O(log N) complexity where N is number of players.
     */
    public void updateScore(String playerId, double score) {
        zSetOps.add(LEADERBOARD_KEY, playerId, score);
    }

    /**
     * Increment a player's score by delta.
     * Useful for adding points from game events.
     */
    public Double incrementScore(String playerId, double delta) {
        return zSetOps.incrementScore(LEADERBOARD_KEY, playerId, delta);
    }

    /**
     * Get a player's rank (1-indexed, where 1 is highest score).
     * O(log N) complexity.
     */
    public Long getRank(String playerId) {
        Long rank = zSetOps.reverseRank(LEADERBOARD_KEY, playerId);
        return (rank != null) ? rank + 1 : null; // Convert to 1-indexed
    }

    /**
     * Get a player's current score.
     */
    public Double getScore(String playerId) {
        return zSetOps.score(LEADERBOARD_KEY, playerId);
    }
}
```

```
/**  
 * Get top N players with their scores.  
 * O(log N + M) where M is the limit.  
 */  
public List<LeaderboardEntry> getTopPlayers(int limit) {  
    Set<ZSetOperations.TypedTuple<String>> results =  
        zSetOps.reverseRangeWithScores(LEADERBOARD_KEY, 0, limit - 1);  
  
    if (results == null) {  
        return Collections.emptyList();  
    }  
  
    List<LeaderboardEntry> entries = new ArrayList<>();  
    int rank = 1;  
  
    for (ZSetOperations.TypedTuple<String> tuple : results) {  
        entries.add(new LeaderboardEntry(  
            tuple.getValue(),  
            tuple.getScore(),  
            rank++  
        ));  
    }  
  
    return entries;  
}  
  
/**  
 * Get players near a specific player (contextual view).  
 * Shows the player's neighbors in the ranking.  
 */  
public List<LeaderboardEntry> getPlayerContext(String playerId, int  
contextSize) {  
    Long rank = zSetOps.reverseRank(LEADERBOARD_KEY, playerId);  
  
    if (rank == null) {  
        return Collections.emptyList();  
    }  
  
    long start = Math.max(0, rank - contextSize);  
    long end = rank + contextSize;  
  
    Set<ZSetOperations.TypedTuple<String>> results =  
        zSetOps.reverseRangeWithScores(LEADERBOARD_KEY, start, end);  
  
    if (results == null) {  
        return Collections.emptyList();  
    }  
  
    List<LeaderboardEntry> entries = new ArrayList<>();  
    int currentRank = (int) start + 1;  
  
    for (ZSetOperations.TypedTuple<String> tuple : results) {  
        entries.add(new LeaderboardEntry(  
            tuple.getValue(),
```

```

        tuple.getScore(),
        currentRank++
    ));
}

return entries;
}

/**
 * Get total number of players on the leaderboard.
 */
public Long getTotalPlayers() {
    return zSetOps.size(LEADERBOARD_KEY);
}

/**
 * Get players within a score range.
 */
public List<LeaderboardEntry> getPlayersByScoreRange(double minScore,
double maxScore) {
    Set<ZSetOperations.TypedTuple<String>> results =
        zSetOps.reverseRangeByScoreWithScores(LEADERBOARD_KEY,
minScore, maxScore);

    if (results == null) {
        return Collections.emptyList();
    }

    return results.stream()
        .map(tuple -> new LeaderboardEntry(
            tuple.getValue(),
            tuple.getScore(),
            0)) // Rank not applicable for range query
        .collect(Collectors.toList());
}
}

/**
 * Leaderboard entry containing player info, score, and rank.
 */
@Data
@AllArgsConstructor
public class LeaderboardEntry {
    private String playerId;
    private Double score;
    private int rank;
}

```

Sharding for Millions of Players:

Shard by score range:
- Shard 1: scores 0-1000

```

- Shard 2: scores 1001-2000
- Shard 3: scores 2001-3000
...
Or by player_id hash:
shard = hash(player_id) % num_shards

```

39. Hotel Reservation System (Detailed)

Covered comprehensively in Problem #15 with double-booking prevention strategies.

40. Multilingual Database Schema

Problem Overview

Design database schema to support multiple languages for product catalogs, content, etc.

Approaches

1. Separate Translation Tables:

```

products (
    id BIGINT PRIMARY KEY,
    sku VARCHAR(50),
    price DECIMAL(10,2)
);

product_translations (
    product_id BIGINT REFERENCES products(id),
    language_code VARCHAR(5),
    name VARCHAR(255),
    description TEXT,
    PRIMARY KEY (product_id, language_code)
);

-- Query
SELECT p.*, pt.name, pt.description
FROM products p
JOIN product_translations pt ON p.id = pt.product_id
WHERE pt.language_code = 'es';

```

2. JSONB Columns:

```

products (
    id BIGINT PRIMARY KEY,
    sku VARCHAR(50),
    price DECIMAL(10,2),

```

```
    names JSONB, -- {"en": "Laptop", "es": "Portátil", "fr": "Ordinateur portable"}  
    descriptions JSONB  
);  
  
-- Query  
SELECT id, names->>'es' as name  
FROM products;
```

3. Separate Tables per Language (not recommended):

```
products_en (...)  
products_es (...)  
products_fr (...)  
-- Maintenance nightmare
```

Best Practice: Separate translation tables for normalized data, JSONB for simpler use cases.

41. System Improvement Analysis

Problem Overview

Given an existing badly designed system, identify flaws and propose improvements with technical justifications.

Analysis Framework

1. Identify Issues:

- Performance bottlenecks
- Scalability limitations
- Single points of failure
- Security vulnerabilities
- Poor data modeling
- Lack of monitoring

2. Propose Solutions:

- Caching strategies
- Database indexing
- Load balancing
- Replication/sharding
- API rate limiting
- Observability tools

Example Analysis:

Current System Issues:

1. Single PostgreSQL instance → Add read replicas + connection pooling
2. No caching → Add Redis cache layer (80% hit rate target)
3. Synchronous payment processing → Async queue (Kafka)
4. No rate limiting → Implement token bucket algorithm
5. Large table scans → Add composite indexes on query patterns

Expected Improvements:


- Latency: 2000ms → 200ms (p95)
- Throughput: 100 req/s → 1000 req/s
- Availability: 99% → 99.9%

42. Multi-Property Hotel Management

Problem Overview

Design a platform for managing multiple hotel properties with centralized inventory, bookings, and analytics.

Architecture



Data Isolation:

```

-- Shared schema with tenant_id
bookings (
    id UUID PRIMARY KEY,
    property_id BIGINT NOT NULL,
    room_id BIGINT,
  
```

```

    ...
);

-- Row-level security
CREATE POLICY property_isolation ON bookings
FOR ALL
USING (property_id = current_setting('app.current_property_id')::bigint);


```

43. Scalable Android System

Problem Overview

Design a robust Android architecture with caching, offline support, and state management.

Architecture Layers



Caching Strategy:

```

class Repository {
    private val api: ApiService,
    private val cache: LocalCache
) {
    suspend fun getData(): Result<Data> {
        // Try cache first
        val cached = cache.get()
        if (cached != null && !cached.isStale()) {
            return Result.Success(cached)
        }

        // Fetch from network
        return try {
            val fresh = api.fetchData()
            cache.save(fresh)
            Result.Success(fresh)
        } catch (e: Exception) {
            // Fallback to stale cache
            cached?.let { Result.Success(it) }
        }
    }
}

```

```
        ?: Result.Error(e)
    }
}
```

44. Flight Inventory with Metered APIs

Covered comprehensively in Problem #25 with cost optimization and rate limiting strategies.

45. Store Inventory Management

Problem Overview

Design an inventory management system for retail stores with real-time stock tracking, reorder alerts, and multi-location support.

Features

- Stock level tracking
- Automatic reorder points
- Transfer between locations
- Audit trail for stock movements
- Low stock alerts

Schema:

```
products (
    id BIGSERIAL PRIMARY KEY,
    sku VARCHAR(50) UNIQUE,
    name VARCHAR(255),
    category VARCHAR(100),
    reorder_point INT,
    reorder_quantity INT
);

locations (
    id BIGSERIAL PRIMARY KEY,
    name VARCHAR(255),
    address TEXT,
    type VARCHAR(20) -- warehouse, store
);

inventory (
    id BIGSERIAL PRIMARY KEY,
    product_id BIGINT REFERENCES products(id),
    location_id BIGINT REFERENCES locations(id),
    quantity INT NOT NULL DEFAULT 0,
    last_updated TIMESTAMP DEFAULT NOW(),
```

```

        UNIQUE(product_id, location_id)
    );

stock_movements (
    id BIGSERIAL PRIMARY KEY,
    product_id BIGINT,
    location_id BIGINT,
    movement_type VARCHAR(20), -- in, out, transfer, adjustment
    quantity INT,
    reference_id VARCHAR(100), -- order_id, transfer_id, etc.
    created_at TIMESTAMP DEFAULT NOW()
);

-- Trigger for low stock alerts
CREATE OR REPLACE FUNCTION check_reorder_point()
RETURNS TRIGGER AS $$ 
BEGIN
    IF NEW.quantity <= (
        SELECT reorder_point FROM products WHERE id = NEW.product_id
    ) THEN
        INSERT INTO reorder_alerts (product_id, location_id, current_quantity)
        VALUES (NEW.product_id, NEW.location_id, NEW.quantity);
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER inventory_reorder_check
AFTER UPDATE OF quantity ON inventory
FOR EACH ROW
EXECUTE FUNCTION check_reorder_point();

```

Inventory Operations:

```

import org.springframework.stereotype.Service;
import org.springframework.transaction.annotation.Transactional;
import org.springframework.jdbc.core.JdbcTemplate;
import java.util.*;

/**
 * Retail store inventory management service.
 * Handles stock adjustments, transfers, and reorder alerts.
 */
@Service
public class InventoryService {

    private final JdbcTemplate db;

    @Autowired
    public InventoryService(JdbcTemplate db) {
        this.db = db;
    }
}

```

```
/**  
 * Adjust stock quantity with full audit trail.  
 * Supports in, out, transfer, and adjustment movements.  
 */  
@Transactional  
public void adjustStock(Long productId, Long locationId, int quantity,  
                        MovementType movementType, String referenceId)  
{  
  
    // Update inventory using UPSERT  
    String upsertQuery = """"  
        INSERT INTO inventory (product_id, location_id, quantity)  
        VALUES (?, ?, ?)  
        ON CONFLICT (product_id, location_id)  
        DO UPDATE SET  
            quantity = inventory.quantity + EXCLUDED.quantity,  
            last_updated = NOW()  
        """;  
  
    db.update(upsertQuery, productId, locationId, quantity);  
  
    // Record movement for audit trail  
    String movementQuery = """"  
        INSERT INTO stock_movements  
        (product_id, location_id, movement_type, quantity,  
        reference_id)  
        VALUES (?, ?, ?, ?, ?, ?)  
        """;  
  
    db.update(movementQuery, productId, locationId,  
              movementType.name().toLowerCase(), quantity, referenceId);  
}  
  
/**  
 * Transfer stock between locations atomically.  
 */  
@Transactional  
public void transferStock(Long productId, Long fromLocation,  
                         Long toLocation, int quantity) {  
    // Generate transfer reference ID  
    String transferId = "TXF-" +  
    UUID.randomUUID().toString().substring(0, 8);  
  
    // Verify source has sufficient stock  
    Integer currentStock = db.queryForObject(  
        "SELECT quantity FROM inventory WHERE product_id = ? AND  
        location_id = ?",
        Integer.class, productId, fromLocation);  
  
    if (currentStock == null || currentStock < quantity) {
        throw new InsufficientStockException(
            "Insufficient stock at source location. Available: " +
            (currentStock != null ? currentStock : 0));
    }
}
```

```
}

    // Deduct from source location
    adjustStock(productId, fromLocation, -quantity,
        MovementType.TRANSFER, transferId);

    // Add to destination location
    adjustStock(productId, toLocation, quantity,
        MovementType.TRANSFER, transferId);
}

/***
 * Get items below reorder point, optionally filtered by location.
 */
public List<LowStockItem> getLowStockItems(Long locationId) {
    String query = """
        SELECT p.id, p.sku, p.name, p.reorder_point,
    p.reorder_quantity,
                i.quantity, i.location_id, l.name as location_name
        FROM products p
        JOIN inventory i ON p.id = i.product_id
        JOIN locations l ON i.location_id = l.id
        WHERE i.quantity <= p.reorder_point
    """;

    if (locationId != null) {
        query += " AND i.location_id = ?";
        return db.query(query, new LowStockItemRowMapper(),
locationId);
    }

    return db.query(query, new LowStockItemRowMapper());
}

/***
 * Get stock level for a product at a specific location.
 */
public Optional<Integer> getStockLevel(Long productId, Long
locationId) {
    String query = """
        SELECT quantity FROM inventory
        WHERE product_id = ? AND location_id = ?
    """;

    try {
        Integer quantity = db.queryForObject(query, Integer.class,
            productId, locationId);
        return Optional.ofNullable(quantity);
    } catch (EmptyResultDataAccessException e) {
        return Optional.empty();
    }
}

/***
```

```
* Get stock movement history for a product.  
*/  
public List<StockMovement> getMovementHistory(Long productId,  
                                              Long locationId,  
                                              int limit) {  
    String query = """  
        SELECT * FROM stock_movements  
        WHERE product_id = ? AND location_id = ?  
        ORDER BY created_at DESC  
        LIMIT ?  
    """;  
  
    return db.query(query, new StockMovementRowMapper(),  
                   productId, locationId, limit);  
}  
}  
  
/**  
 * Types of inventory movements.  
 */  
public enum MovementType {  
    IN,           // Stock received  
    OUT,          // Stock sold/shipped  
    TRANSFER,     // Inter-location transfer  
    ADJUSTMENT   // Manual correction  
}  
  
/**  
 * Low stock item with reorder information.  
 */  
@Data  
public class LowStockItem {  
    private Long productId;  
    private String sku;  
    private String name;  
    private int currentQuantity;  
    private int reorderPoint;  
    private int reorderQuantity;  
    private Long locationId;  
    private String locationName;  
}  
  
/**  
 * Stock movement audit record.  
 */  
@Data  
public class StockMovement {  
    private Long id;  
    private Long productId;  
    private Long locationId;  
    private MovementType movementType;  
    private int quantity;  
    private String referenceId;  
    private Instant createdAt;
```

```
}

/**
 * Exception for insufficient stock scenarios.
 */
public class InsufficientStockException extends RuntimeException {
    public InsufficientStockException(String message) {
        super(message);
    }
}
```

Summary of All 45 Solutions

This document provides comprehensive system design solutions covering:

Core Systems (1-15):

- Streaming platforms, search systems, file storage
- Booking systems with concurrency control
- Distributed scheduling, payment gateways

Concepts & Patterns (16-23):

- Caching strategies, sharding, consistency models
- Rate limiting, top-K algorithms

Specialized Systems (24-45):

- Reconciliation, flight inventory
- Distributed KV stores, message brokers
- Real-time chat, leaderboards
- Multi-tenant platforms

Key Themes:

- Horizontal scalability via partitioning/sharding
- High availability through replication
- Low latency via multi-level caching
- Strong consistency where needed (transactions)
- Eventual consistency where acceptable (performance)
- Observability and monitoring throughout

Each solution includes: Architecture diagrams Database schemas Implementation code
Trade-off analysis Scalability considerations

End of Complete System Design Solutions (1-45)