

System Design Solutions: Problems 26-45

Complete architectural solutions with diagrams, code examples, and implementation details.

26. Distributed Key-Value Store

Problem Overview

Design a distributed key-value data store like Redis or DynamoDB that provides high availability, horizontal scalability, and tunable consistency guarantees.

Back-of-the-Envelope Estimation

- **Operations/sec:** 1 million ops/sec (70% reads, 30% writes)
- **Data size:** 1TB total data
- **Key size:** Average 50 bytes
- **Value size:** Average 1KB
- **Total keys:** ~1 billion keys
- **Nodes:** 10 nodes initially, scalable to 100+
- **Replication factor:** 3 (for durability)
- **Storage per node:** 100GB with replication

Functional Requirements

- **FR1:** Put(key, value) - Store key-value pair
- **FR2:** Get(key) - Retrieve value by key
- **FR3:** Delete(key) - Remove key-value pair
- **FR4:** Support TTL (Time To Live) for automatic expiration
- **FR5:** Atomic operations (increment, compare-and-swap)
- **FR6:** Range queries (scan keys by prefix)

Non-Functional Requirements

- **Scalability:** Horizontal scaling to 100+ nodes
- **Availability:** 99.99% uptime (4 nines)
- **Latency:** <1ms for reads, <5ms for writes (p99)
- **Consistency:** Tunable (strong, eventual, or read-your-writes)
- **Durability:** No data loss (replicated to N nodes)
- **Partition Tolerance:** Continue operating during network partitions

High-Level Architecture

Components:


- **Client Library:** Smart client with routing logic
- **Coordinator Nodes:** Handle requests, route to correct partition
- **Storage Nodes:** Store actual data with LSM trees


- **Replication Manager:** Ensure N replicas per key
- **Membership Service:** Track node health (gossip protocol)
- **Partition Manager:** Consistent hashing for data distribution

Data Storage Choices

| Component | Technology | Justification |
|----------------|--------------------|---|
| Storage Engine | LSM Tree (RocksDB) | High write throughput, efficient compaction |
| Replication | Multi-master | High availability, low latency writes |
| Partitioning | Consistent Hashing | Even distribution, minimal data movement |
| Membership | Gossip Protocol | Decentralized, fault-tolerant |

Architecture Diagram:





Replication Strategy (N=3):

Key "user:123" → hash → position 45°

Primary replicas (clockwise):

1. Node 2 (45°) ← Primary/Leader
2. Node 3 (90°) ← Replica
3. Node 1 (135°) ← Replica

Write Path (Quorum W=2):

1. Client → hash(key) → find replicas [N2, N3, N1]
2. Send write to all 3 replicas in parallel
3. Wait for W=2 acknowledgments
4. Return success to client
5. Async: 3rd replica eventually catches up

Read Path (Quorum R=2):

1. Client → hash(key) → find replicas [N2, N3, N1]
2. Send read to R=2 replicas
3. If values match → return
4. If values differ → read repair (return latest)

Implementation:

```

import hashlib
import bisect
from typing import Any, List, Optional, Tuple
from datetime import datetime, timedelta
import asyncio

class ConsistentHashRing:
    """Consistent hashing ring for data distribution."""

    def __init__(self, nodes: List[str], virtual_nodes: int = 150):
        self.virtual_nodes = virtual_nodes
        self.ring = [] # Sorted list of (hash, node) tuples
        self.nodes = set()
  
```

```
for node in nodes:
    self.add_node(node)

def _hash(self, key: str) -> int:
    """Hash key to position on ring (0-2^32)."""
    return int(hashlib.md5(key.encode()).hexdigest(), 16)

def add_node(self, node: str):
    """Add a node with virtual nodes to the ring."""
    self.nodes.add(node)
    for i in range(self.virtual_nodes):
        virtual_key = f"{node}:{i}"
        hash_value = self._hash(virtual_key)
        bisect.insort(self.ring, (hash_value, node))

def remove_node(self, node: str):
    """Remove a node and its virtual nodes."""
    self.nodes.discard(node)
    self.ring = [(h, n) for h, n in self.ring if n != node]

def get_nodes(self, key: str, count: int = 3) -> List[str]:
    """Get N responsible nodes for a key (clockwise)."""
    if not self.ring:
        return []

    hash_value = self._hash(key)

    # Find position in ring
    idx = bisect.bisect(self.ring, (hash_value, ''))

    # Collect unique nodes clockwise
    nodes = []
    seen = set()

    for i in range(len(self.ring)):
        pos = (idx + i) % len(self.ring)
        node = self.ring[pos][1]

        if node not in seen:
            nodes.append(node)
            seen.add(node)

        if len(nodes) == count:
            break

    return nodes

class DistributedKVStore:
    """Main distributed key-value store."""

    def __init__(self, node_id: str, nodes: List[str],
                 replication_factor: int = 3):
        self.node_id = node_id
```

```
self.replication_factor = replication_factor
self.hash_ring = ConsistentHashRing(nodes)

# Local storage (LSM tree abstraction)
self.storage = LSMTTreeStorage()

# Replication configuration
self.write_quorum = 2 # W
self.read_quorum = 2 # R
# Guarantee: R + W > N ensures consistency

# Node connections
self.node_connections = {n: NodeConnection(n) for n in nodes}

async def put(self, key: str, value: Any, ttl: Optional[int] = None) -> bool:
    """
    Store key-value pair with optional TTL.

    Args:
        key: Key to store
        value: Value to store
        ttl: Time to live in seconds (optional)

    Returns:
        True if write succeeded (quorum reached)
    """
    # Find replica nodes
    replicas = self.hash_ring.get_nodes(key, self.replication_factor)

    # Create versioned value
    versioned_value = VersionedValue(
        value=value,
        timestamp=datetime.now(),
        version=self._generate_version(),
        ttl=ttl
    )

    # Write to all replicas in parallel
    write_tasks = []
    for replica in replicas:
        if replica == self.node_id:
            # Local write
            task = self._local_put(key, versioned_value)
        else:
            # Remote write
            task = self._remote_put(replica, key, versioned_value)

        write_tasks.append(task)

    # Wait for quorum
    results = await asyncio.gather(*write_tasks,
return_exceptions=True)
```

```
# Count successful writes
success_count = sum(1 for r in results if r is True)

if success_count >= self.write_quorum:
    # Quorum reached
    return True
else:
    # Quorum not reached - rollback or retry
    await self._handle_write_failure(key, replicas)
    return False

async def get(self, key: str) -> Optional[Any]:
    """
    Retrieve value by key with quorum read.

    Returns:
        Value if found, None otherwise
    """

    # Find replica nodes
    replicas = self.hash_ring.get_nodes(key, self.replication_factor)

    # Read from R replicas in parallel
    read_tasks = []
    for i, replica in enumerate(replicas[:self.read_quorum]):
        if replica == self.node_id:
            task = self._local_get(key)
        else:
            task = self._remote_get(replica, key)

        read_tasks.append(task)

    # Wait for quorum responses
    results = await asyncio.gather(*read_tasks,
return_exceptions=True)

    # Filter valid results
    valid_results = [r for r in results if isinstance(r,
VersionedValue)]

    if not valid_results:
        return None

    # Resolve conflicts (last-write-wins)
    latest = max(valid_results, key=lambda v: v.timestamp)

    # Check TTL
    if latest.ttl and latest.is_expired():
        await self.delete(key)
        return None

    # Read repair if needed
    if len(set(r.version for r in valid_results)) > 1:
        await self._read_repair(key, latest, replicas)
```

```
        return latest.value

    async def delete(self, key: str) -> bool:
        """Delete key-value pair."""
        # Use tombstone for eventual consistency
        tombstone = VersionedValue(
            value=None,
            timestamp=datetime.now(),
            version=self._generate_version(),
            is_tombstone=True
        )

        return await self.put(key, tombstone)

    async def _local_put(self, key: str, value: VersionedValue) -> bool:
        """Write to local storage."""
        try:
            self.storage.put(key, value)
            return True
        except Exception as e:
            print(f"Local write failed: {e}")
            return False

    async def _local_get(self, key: str) -> Optional[VersionedValue]:
        """Read from local storage."""
        return self.storage.get(key)

    async def _remote_put(self, node: str, key: str,
                         value: VersionedValue) -> bool:
        """Write to remote node."""
        try:
            conn = self.node_connections[node]
            return await conn.put(key, value)
        except Exception as e:
            print(f"Remote write to {node} failed: {e}")
            return False

    async def _remote_get(self, node: str, key: str) ->
Optional[VersionedValue]:
        """Read from remote node."""
        try:
            conn = self.node_connections[node]
            return await conn.get(key)
        except Exception as e:
            print(f"Remote read from {node} failed: {e}")
            return None

    async def _read_repair(self, key: str, correct_value: VersionedValue,
                          replicas: List[str]):
        """Repair stale replicas with correct value."""
        for replica in replicas:
            if replica == self.node_id:
                continue
```

```
# Write correct value to replica
await self._remote_put(replica, key, correct_value)

def _generate_version(self) -> str:
    """Generate version using vector clock or timestamp."""
    return f"{self.node_id}:{datetime.now().timestamp()}"

class VersionedValue:
    """Value with version and metadata."""

    def __init__(self, value: Any, timestamp: datetime, version: str,
                 ttl: Optional[int] = None, is_tombstone: bool = False):
        self.value = value
        self.timestamp = timestamp
        self.version = version
        self.ttl = ttl
        self.is_tombstone = is_tombstone

    def is_expired(self) -> bool:
        """Check if value has expired."""
        if not self.ttl:
            return False

        expiry = self.timestamp + timedelta(seconds=self.ttl)
        return datetime.now() > expiry

class LSMTreeStorage:
    """LSM Tree storage engine (simplified)."""

    def __init__(self):
        self.memtable = {} # In-memory write buffer
        self.sstables = [] # Sorted string tables on disk
        self.memtable_size_limit = 1000

    def put(self, key: str, value: VersionedValue):
        """Write to memtable."""
        self.memtable[key] = value

        # Flush to disk if memtable is full
        if len(self.memtable) >= self.memtable_size_limit:
            self._flush_memtable()

    def get(self, key: str) -> Optional[VersionedValue]:
        """Read from memtable, then sstables."""
        # Check memtable first
        if key in self.memtable:
            return self.memtable[key]

        # Check sstables (newest first)
        for sstable in reversed(self.sstables):
            value = sstable.get(key)
            if value:
```

```
        return value

    return None

def _flush_memtable(self):
    """Flush memtable to disk as SSTable."""
    # Sort by key
    sorted_items = sorted(self.memtable.items())

    # Create new SSTable
    sstable = SSTable(sorted_items)
    self.sstables.append(sstable)

    # Clear memtable
    self.memtable = {}

    # Trigger compaction if needed
    if len(self.sstables) > 5:
        self._compact_sstables()

def _compact_sstables(self):
    """Merge and compact SSTables."""
    # Simplified: merge all SSTables
    all_items = {}

    for sstable in self.sstables:
        for key, value in sstable.items():
            # Keep latest version
            if key not in all_items or value.timestamp >
all_items[key].timestamp:
                all_items[key] = value

    # Create new compacted SSTable
    sorted_items = sorted(all_items.items())
    compacted = SSTable(sorted_items)

    # Replace old SSTables
    self.sstables = [compacted]

class SSTable:
    """Sorted String Table (on-disk storage)."""

    def __init__(self, items: List[Tuple[str, VersionedValue]]):
        self.items = dict(items)
        # In real implementation, would write to disk with bloom filter

    def get(self, key: str) -> Optional[VersionedValue]:
        return self.items.get(key)

class NodeConnection:
    """Connection to remote node."""
```

```

def __init__(self, node_address: str):
    self.address = node_address

async def put(self, key: str, value: VersionedValue) -> bool:
    # RPC call to remote node
    # Implementation would use gRPC, HTTP, or custom protocol
    pass

async def get(self, key: str) -> Optional[VersionedValue]:
    # RPC call to remote node
    pass

```

Gossip Protocol for Membership:

```

import random
import time

class GossipMembership:
    """Gossip protocol for node membership and failure detection."""

    def __init__(self, node_id: str, seed_nodes: List[str]):
        self.node_id = node_id
        self.nodes = {} # node_id -> NodeState
        self.heartbeat_interval = 1 # seconds
        self.gossip_fanout = 3 # gossip to N random nodes

        # Initialize with seed nodes
        for node in seed_nodes:
            self.nodes[node] = NodeState(node, is_alive=True)

        self.nodes[node_id] = NodeState(node_id, is_alive=True)

    async def start_gossip(self):
        """Start gossip protocol."""
        while True:
            await asyncio.sleep(self.heartbeat_interval)

            # Increment own heartbeat
            self.nodes[self.node_id].increment_heartbeat()

            # Select random nodes to gossip
            gossip_targets = self._select_gossip_targets()

            # Send gossip to targets
            for target in gossip_targets:
                await self._send_gossip(target)

            # Detect failures
            self._detect_failures()

    def _select_gossip_targets(self) -> List[str]:
        """Select random nodes for gossip."""

```

```
alive_nodes = [n for n, s in self.nodes.items()
               if s.is_alive and n != self.node_id]

count = min(self.gossip_fanout, len(alive_nodes))
return random.sample(alive_nodes, count)

async def _send_gossip(self, target: str):
    """Send gossip message to target node."""
    message = GossipMessage(
        sender=self.node_id,
        node_states=self.nodes
    )

    # Send via network (RPC)
    # await rpc_call(target, message)
    pass

def receive_gossip(self, message: GossipMessage):
    """Process received gossip message."""
    for node_id, remote_state in message.node_states.items():
        if node_id not in self.nodes:
            # New node discovered
            self.nodes[node_id] = remote_state
        else:
            # Merge states (keep higher heartbeat)
            local_state = self.nodes[node_id]
            if remote_state.heartbeat > local_state.heartbeat:
                self.nodes[node_id] = remote_state

def _detect_failures(self):
    """Detect failed nodes based on heartbeat timeouts."""
    now = time.time()
    timeout = 10 # seconds

    for node_id, state in self.nodes.items():
        if node_id == self.node_id:
            continue

        if state.is_alive and (now - state.last_update) > timeout:
            # Node suspected as failed
            state.is_alive = False
            print(f"Node {node_id} marked as failed")

            # Trigger rebalancing
            self._handle_node_failure(node_id)

def _handle_node_failure(self, failed_node: str):
    """Handle node failure - rebalance data."""
    # Remove from hash ring
    # Redistribute data to other nodes
    pass

class NodeState:
```

```
"""State of a node in the cluster."""

def __init__(self, node_id: str, is_alive: bool = True):
    self.node_id = node_id
    self.is_alive = is_alive
    self.heartbeat = 0
    self.last_update = time.time()

def increment_heartbeat(self):
    """Increment heartbeat counter."""
    self.heartbeat += 1
    self.last_update = time.time()
```

Trade-offs & Assumptions

- **Consistency vs Availability:** Tunable quorum ($R+W>N$) allows choosing between strong consistency and high availability
 - **CAP Theorem:** System is AP (Available + Partition Tolerant) by default, CP with strong quorum
 - **Replication Factor:** $N=3$ balances durability and storage cost
 - **Virtual Nodes:** 150 per physical node ensures even distribution
 - **Assumption:** Network partitions are rare; gossip protocol detects failures within 10 seconds
-

27. Movie Seat Booking System

Problem Overview

Design a movie ticket booking system handling concurrent seat reservations with payment integration, preventing double bookings, and managing the complete flow from seat selection to confirmation.

Back-of-the-Envelope Estimation

- **Theaters:** 500 theaters
- **Screens per theater:** 10 screens
- **Shows per screen/day:** 5 shows
- **Seats per show:** 200 seats
- **Total seats/day:** $500 \times 10 \times 5 \times 200 = 5 \text{ million seats}$
- **Booking rate:** 20% occupancy = 1 million bookings/day
- **Peak bookings/sec:** $1M / 86400 \times 10$ (peak factor) = ~115 bookings/sec
- **Concurrent users:** 10K users simultaneously booking

Functional Requirements


- **FR1:** Browse movies, theaters, and showtimes
- **FR2:** View seat map with availability in real-time
- **FR3:** Select seats and hold temporarily (10 minutes)
- **FR4:** Process payment
- **FR5:** Generate and send ticket confirmation
- **FR6:** Cancel bookings with refund

- **FR7:** Handle booking failures and rollbacks

Non-Functional Requirements

- **Scalability:** Handle 115 bookings/sec peak load
- **Availability:** 99.9% uptime
- **Latency:** <500ms for seat selection, <3s for payment
- **Consistency:** Strong consistency (no double bookings)
- **Atomicity:** Seat reservation + payment atomic

High-Level Architecture



```

FOR UPDATE NOWAIT; -- Acquire row-level lock

IF all seats available:
    UPDATE seats SET status = 'held',
        held_by = session_id,
        held_until = NOW() + INTERVAL '10 minutes'
    WHERE ...
ELSE:
    ROLLBACK -- Seats taken
COMMIT

```

3. User Enters Payment

↳ 10 min timer countdown

4. Process Payment

BEGIN TRANSACTION

-- Verify seats still held by this session

SELECT * FROM seats

WHERE show_id = ? AND seat_numbers IN (?)

AND held_by = session_id

FOR UPDATE;

-- Process payment

payment_result = payment_gateway.charge()

IF payment successful:

INSERT INTO bookings (...)

UPDATE seats SET status = 'booked'

COMMIT

ELSE:

UPDATE seats SET status = 'available'

ROLLBACK

END TRANSACTION

5. Generate Ticket

↳ QR code, PDF, email

Double Booking Prevention Strategy:

Layer 1: Application-Level Lock

- Redis distributed lock
- Prevents concurrent hold attempts

Layer 2: Database Row Lock

- SELECT FOR UPDATE NOWAIT
- Database guarantees exclusivity

Layer 3: Unique Constraint

- UNIQUE(show_id, seat_number)
- Final safety net

Database Schema:

```
-- Movies and shows
movies (
    id BIGSERIAL PRIMARY KEY,
    title VARCHAR(255) NOT NULL,
    duration INT NOT NULL, -- minutes
    genre VARCHAR(50),
    rating VARCHAR(10)
);

theaters (
    id BIGSERIAL PRIMARY KEY,
    name VARCHAR(255) NOT NULL,
    location VARCHAR(255),
    total_screens INT
);

screens (
    id BIGSERIAL PRIMARY KEY,
    theater_id BIGINT REFERENCES theaters(id),
    screen_number INT,
    total_seats INT,
    seat_layout JSONB -- {"rows": 10, "cols": 20}
);

shows (
    id BIGSERIAL PRIMARY KEY,
    movie_id BIGINT REFERENCES movies(id),
    screen_id BIGINT REFERENCES screens(id),
    show_time TIMESTAMP NOT NULL,
    price_base DECIMAL(10,2),
    status VARCHAR(20) DEFAULT 'scheduled'
);

-- Seats with locking
seats (
    id BIGSERIAL PRIMARY KEY,
    show_id BIGINT REFERENCES shows(id),
    seat_row VARCHAR(5) NOT NULL,
    seat_number INT NOT NULL,
    seat_type VARCHAR(20), -- regular, premium, vip
    price DECIMAL(10,2) NOT NULL,
    status VARCHAR(20) NOT NULL DEFAULT 'available',
    -- available, held, booked, blocked
    held_by VARCHAR(100), -- session_id
    held_until TIMESTAMP,
    UNIQUE(show_id, seat_row, seat_number)
);

CREATE INDEX idx_seats_show_status ON seats(show_id, status);
CREATE INDEX idx_seats_held_until ON seats(held_until)
```

```

WHERE status = 'held';

-- Bookings
bookings (
    id UUID PRIMARY KEY,
    user_id BIGINT,
    show_id BIGINT REFERENCES shows(id),
    seat_ids BIGINT[],
    total_amount DECIMAL(10,2),
    payment_id VARCHAR(100),
    status VARCHAR(20) NOT NULL,
        -- pending, confirmed, cancelled, refunded
    booking_time TIMESTAMP DEFAULT NOW(),
    confirmation_code VARCHAR(20) UNIQUE
);

-- Payments
payments (
    id BIGSERIAL PRIMARY KEY,
    booking_id UUID REFERENCES bookings(id),
    amount DECIMAL(10,2) NOT NULL,
    payment_method VARCHAR(20),
    gateway_transaction_id VARCHAR(100),
    status VARCHAR(20) NOT NULL,
        -- initiated, success, failed, refunded
    created_at TIMESTAMP DEFAULT NOW()
);

```

Booking Service Implementation:

```

class BookingService:
    """Handle movie ticket bookings with concurrency control."""

    def __init__(self, db, redis, payment_service):
        self.db = db
        self.redis = redis
        self.payment = payment_service
        self.hold_duration = 600 # 10 minutes in seconds

    @asyncio.coroutine
    def hold_seats(self, show_id: int, seat_ids: List[int],
                   session_id: str) -> Dict:
        """
        Hold seats for a user temporarily.

        Uses distributed lock + database row lock for double booking
        prevention.
        """
        # Generate lock key
        lock_key = f"booking:show:{show_id}:seats:{','.join(map(str,
seat_ids))}"

```

Acquire distributed lock (prevents concurrent attempts)

```
lock_acquired = await self._acquire_distributed_lock(lock_key,
session_id)

if not lock_acquired:
    return {
        'success': False,
        'error': 'CONCURRENT_BOOKING',
        'message': 'Someone else is booking these seats'
    }

try:
    # Database transaction with row-level locking
    async with self.db.transaction(isolation='Serializable'):
        # Lock seats (NOWAIT fails immediately if locked)
        query = """
SELECT id, status, held_by, held_until
FROM seats
WHERE show_id = $1 AND id = ANY($2)
FOR UPDATE NOWAIT
"""

        try:
            seats = await self.db.fetch(query, show_id, seat_ids)
        except asynccpg.exceptions.LockNotFoundError:
            return {
                'success': False,
                'error': 'SEATS_LOCKED',
                'message': 'Seats are being booked by another
user'
            }

        # Validate all seats are available
        for seat in seats:
            if seat['status'] == 'booked':
                return {
                    'success': False,
                    'error': 'SEATS_ALREADY_BOOKED',
                    'message': f"Seat {seat['id']} is already
booked"
                }

            if seat['status'] == 'held' and seat['held_by'] != session_id:
                # Check if hold expired
                if seat['held_until'] > datetime.now():
                    return {
                        'success': False,
                        'error': 'SEATS_HELD',
                        'message': f"Seat {seat['id']} is held by
another user"
                    }

                # Hold seats
                held_until = datetime.now() +
```

```
timedelta(seconds=self.hold_duration)

        update_query = """
        UPDATE seats
        SET status = 'held',
            held_by = $1,
            held_until = $2
        WHERE show_id = $3 AND id = ANY($4)
        """

        await self.db.execute(update_query, session_id,
hold_until,
                           show_id, seat_ids)

        # Commit transaction
        # Locks released automatically on commit

        return {
            'success': True,
            'held_until': held_until.isoformat(),
            'hold_duration': self.hold_duration
        }

    finally:
        # Release distributed lock
        await self._release_distributed_lock(lock_key, session_id)

async def confirm_booking(self, show_id: int, seat_ids: List[int],
                           session_id: str, user_id: int,
                           payment_details: Dict) -> Dict:
    """
    Confirm booking with payment processing.

    Implements Saga pattern for distributed transaction.
    """
    booking_id = str(uuid.uuid4())

    try:
        # Step 1: Verify seats still held by this session
        async with self.db.transaction(isolation='Serializable'):
            verify_query = """
            SELECT id, price, status, held_by
            FROM seats
            WHERE show_id = $1 AND id = ANY($2)
            FOR UPDATE
            """

            seats = await self.db.fetch(verify_query, show_id,
seat_ids)

            # Validate ownership and status
            for seat in seats:
                if seat['status'] != 'held' or seat['held_by'] !=

session_id:
```

```
        raise BookingError('SEATS_NOT_HELD',
                            'Seats no longer held by this
session')

    total_amount = sum(seat['price'] for seat in seats)

    # Create booking record (pending)
    booking_query = """
    INSERT INTO bookings
    (id, user_id, show_id, seat_ids, total_amount, status)
    VALUES ($1, $2, $3, $4, $5, 'pending')
    RETURNING confirmation_code
"""

    result = await self.db.fetchrow(
        booking_query, booking_id, user_id, show_id,
        seat_ids, total_amount
    )

    confirmation_code = result['confirmation_code']

# Step 2: Process payment (external service)
try:
    payment_result = await self.payment.charge(
        amount=total_amount,
        currency='USD',
        payment_method=payment_details,
        metadata={'booking_id': booking_id}
    )

    if not payment_result['success']:
        raise PaymentError(payment_result['error'])

    payment_id = payment_result['transaction_id']

except PaymentError as e:
    # Payment failed - rollback booking
    await self._rollback_booking(booking_id, show_id,
seat_ids)

    return {
        'success': False,
        'error': 'PAYMENT_FAILED',
        'message': str(e)
    }

# Step 3: Finalize booking
async with self.db.transaction():
    # Update seats to booked
    update_seats_query = """
    UPDATE seats
    SET status = 'booked',
        held_by = NULL,
        held_until = NULL
"""

    await self.db.execute(update_seats_query)
```

```
        WHERE show_id = $1 AND id = ANY($2)
        .....

        await self.db.execute(update_seats_query, show_id,
seat_ids)

        # Update booking status
        update_booking_query = """
        UPDATE bookings
        SET status = 'confirmed',
            payment_id = $1
        WHERE id = $2
        .....

        await self.db.execute(update_booking_query, payment_id,
booking_id)

        # Record payment
        payment_query = """
        INSERT INTO payments
        (booking_id, amount, payment_method,
gateway_transaction_id, status)
        VALUES ($1, $2, $3, $4, 'success')
        .....

        await self.db.execute(
            payment_query, booking_id, total_amount,
            payment_details['method'], payment_id
        )

# Step 4: Generate ticket (async)
asyncio.create_task(
    self._generate_and_send_ticket(booking_id, user_id)
)

return {
    'success': True,
    'booking_id': booking_id,
    'confirmation_code': confirmation_code,
    'total_amount': float(total_amount)
}

except BookingError as e:
    return {
        'success': False,
        'error': e.code,
        'message': e.message
    }

except Exception as e:
    # Unexpected error - attempt rollback
    await self._rollback_booking(booking_id, show_id, seat_ids)
    raise
```

```
async def _rollback_booking(self, booking_id: str, show_id: int,
                           seat_ids: List[int]):
    """Rollback booking in case of failure."""
    async with self.db.transaction():
        # Release seats
        await self.db.execute("""
            UPDATE seats
            SET status = 'available',
                held_by = NULL,
                held_until = NULL
            WHERE show_id = $1 AND id = ANY($2)
        """, show_id, seat_ids)

        # Mark booking as cancelled
        await self.db.execute("""
            UPDATE bookings
            SET status = 'cancelled'
            WHERE id = $1
        """, booking_id)

async def release_expired_holds(self):
    """Background job to release expired seat holds."""
    while True:
        await asyncio.sleep(60) # Run every minute

        async with self.db.transaction():
            # Find expired holds
            expired = await self.db.fetch("""
                SELECT DISTINCT show_id, ARRAY_AGG(id) as seat_ids
                FROM seats
                WHERE status = 'held' AND held_until < NOW()
                GROUP BY show_id
            """)

            for row in expired:
                # Release seats
                await self.db.execute("""
                    UPDATE seats
                    SET status = 'available',
                        held_by = NULL,
                        held_until = NULL
                    WHERE show_id = $1 AND id = ANY($2)
                """, row['show_id'], row['seat_ids'])

                print(f"Released {len(row['seat_ids'])} expired holds
for show {row['show_id']}")

    async def _acquire_distributed_lock(self, key: str, value: str) ->
bool:
    """Acquire distributed lock using Redis."""
    result = await self.redis.set(
        key, value,
        nx=True, # Only set if not exists
        ex=30 # 30 second expiry
```

```

        )
    return result is not None

    async def _release_distributed_lock(self, key: str, value: str):
        """Release distributed lock."""
        # Lua script ensures we only delete if we own the lock
        script = """
            if redis.call("GET", KEYS[1]) == ARGV[1] then
                return redis.call("DEL", KEYS[1])
            else
                return 0
            end
        """
        await self.redis.eval(script, 1, key, value)

```

Trade-offs & Assumptions

- **Pessimistic Locking:** `SELECT FOR UPDATE NOWAIT` prevents concurrent bookings but reduces concurrency
 - **Hold Duration:** 10 minutes balances user experience vs inventory blocking
 - **Saga Pattern:** Compensating transactions handle payment failures
 - **Assumption:** 80% of holds result in successful bookings
 - **Distributed Lock:** Redis lock prevents concurrent hold attempts before DB lock
-

28. E-commerce Top Sellers

Problem Overview

Design a system to track and display top 20-30 selling items in real-time for an e-commerce platform with millions of products and high transaction volume.

Back-of-the-Envelope Estimation

- **Products:** 10 million SKUs
- **Orders/day:** 5 million orders
- **Items per order:** Average 2.5 items
- **Total item purchases/day:** 12.5 million
- **Peak purchases/sec:** $12.5M / 86400 \times 10 = \sim 1,450/\text{sec}$
- **Top sellers update frequency:** Every 5 minutes
- **Categories:** 1,000 categories

Functional Requirements


- **FR1:** Track item purchase counts in real-time
- **FR2:** Calculate top 20-30 sellers globally
- **FR3:** Calculate top sellers per category
- **FR4:** Update rankings every 5 minutes
- **FR5:** Display trending items (velocity-based)



- **FR6:** Historical trending data (daily, weekly, monthly)

Non-Functional Requirements

- **Scalability:** Handle 1,450 purchases/sec
- **Latency:** <100ms for fetching top sellers
- **Accuracy:** 99% accuracy acceptable (approximate counts)
- **Freshness:** Rankings updated every 5 minutes

Architecture





Redis Data Structures:

ZADD top:global:current product:123 5420

ZADD top:category:electronics product:111 2103

ZADD trending:velocity product:456 892 # purchases/hour

Query API:

GET /api/top-sellers?limit=30

→ ZREVRANGE top:global:current 0 29 WITHSCORES

GET /api/top-sellers/category/electronics?limit=20

→ ZREVRANGE top:category:electronics 0 19 WITHSCORES

Stream Processing (Flink):

```

from pyflink.datastream import StreamExecutionEnvironment
from pyflink.datastream.window import TumblingProcessingTimeWindows
from pyflink.common import Time

def process_purchases_stream():
    env = StreamExecutionEnvironment.get_execution_environment()
    env.set_parallelism(3)

    # Source: Kafka
    purchases_stream = env.add_source(
        FlinkKafkaConsumer(
            topics=['purchase-events'],
            deserialization_schema=JsonDeserializationSchema(),
            properties={'bootstrap.servers': 'kafka:9092'}
        )
    )

    # Extract product_id and category
  
```

```

product_counts = purchases_stream \
    .map(lambda x: (x['product_id'], x['category'], 1)) \
    .key_by(lambda x: x[0]) \
    .window(TumblingProcessingTimeWindows.of(Time.minutes(5))) \
    .sum(2) # Sum counts

# Update Redis sorted sets
product_counts.add_sink(RedisTopSellersSink())

env.execute("Top Sellers Stream")

```



```

class RedisTopSellersSink:
    """Sink to update Redis sorted sets."""

    def __init__(self):
        self.redis = redis.Redis(host='redis')

    def invoke(self, value, context):
        product_id, category, count = value

        # Update global top sellers
        self.redis.zincrby('top:global:current', count, f'product:{product_id}')

        # Update category top sellers
        self.redis.zincrby(f'top:category:{category}', count, f'product:{product_id}')

        # Calculate velocity (trending)
        # purchases in last hour / purchases in previous hour
        velocity = self._calculate_velocity(product_id)
        self.redis.zadd('trending:velocity', {f'product:{product_id}': velocity})

    def _calculate_velocity(self, product_id):
        """Calculate trending velocity."""
        current_hour = self.redis.get(f'count:hour:current:{product_id}') or 0
        previous_hour = self.redis.get(f'count:hour:previous:{product_id}') or 1

        return float(current_hour) / float(previous_hour)

```

API Service:

```

from fastapi import FastAPI, Query
from typing import List, Optional

app = FastAPI()

class TopSellersService:

```

```
def __init__(self, redis_client, db):
    self.redis = redis_client
    self.db = db

    async def get_top_sellers(self, limit: int = 30,
                               category: Optional[str] = None) ->
List[Dict]:
        """Get top selling products."""
        # Determine Redis key
        if category:
            key = f'top:category:{category}'
        else:
            key = 'top:global:current'

        # Get from Redis sorted set (highest scores first)
        results = self.redis.zrevrange(key, 0, limit - 1, withscores=True)

        # Extract product IDs
        product_ids = [r[0].decode().split(':')[1] for r in results]

        # Fetch product details from DB
        products = await self.db.fetch("""
            SELECT id, name, price, image_url, category
            FROM products
            WHERE id = ANY($1)
        """, product_ids)

        # Merge with scores
        product_map = {str(p['id']): p for p in products}

        top_sellers = []
        for product_key, score in results:
            product_id = product_key.decode().split(':')[1]

            if product_id in product_map:
                product = product_map[product_id]
                top_sellers.append({
                    'product_id': product_id,
                    'name': product['name'],
                    'price': float(product['price']),
                    'image_url': product['image_url'],
                    'category': product['category'],
                    'purchase_count': int(score),
                    'rank': len(top_sellers) + 1
                })

        return top_sellers

    async def get_trending_products(self, limit: int = 20) -> List[Dict]:
        """Get trending products based on velocity."""
        # Get products with highest velocity
        results = self.redis.zrevrange('trending:velocity', 0, limit - 1,
                                       withscores=True)
```

```

product_ids = [r[0].decode().split(':')[1] for r in results]

products = await self.db.fetch("""
    SELECT id, name, price, image_url
    FROM products
    WHERE id = ANY($1)
""", product_ids)

product_map = {str(p['id']): p for p in products}

trending = []
for product_key, velocity in results:
    product_id = product_key.decode().split(':')[1]

    if product_id in product_map:
        product = product_map[product_id]
        trending.append({
            'product_id': product_id,
            'name': product['name'],
            'price': float(product['price']),
            'image_url': product['image_url'],
            'velocity': float(velocity),
            'trend': 'hot' if velocity > 2 else 'rising'
        })

return trending

@app.get("/api/top-sellers")
async def get_top_sellers(
    limit: int = Query(30, ge=1, le=100),
    category: Optional[str] = None
):
    """Get top selling products."""
    service = TopSellersService(redis, db)
    return await service.get_top_sellers(limit, category)

@app.get("/api/trending")
async def get_trending(limit: int = Query(20, ge=1, le=50)):
    """Get trending products."""
    service = TopSellersService(redis, db)
    return await service.get_trending_products(limit)

```

Trade-offs & Assumptions

- **Approximate Counts:** Redis sorted sets provide fast queries but approximate counts acceptable
- **Update Frequency:** 5-minute windows balance freshness vs computational cost
- **Top K Size:** Tracking top 30 globally, top 20 per category limits memory usage
- **Trending Algorithm:** Velocity-based (current/previous) simple but effective
- **Assumption:** 80% of sales come from 20% of products (Pareto principle)

[Continuing with remaining problems 29–45...]

29. Multi-Datacenter Replication

Problem Overview

Design a strategy for replicating data across multiple datacenters globally while maintaining consistency, minimizing latency, and handling network partitions.

Replication Strategies

1. Master-Slave Replication:

Primary DC (US-East)
↓ Async Replication
Secondary DCs (EU, APAC)

Writes: US-East only
Reads: Any DC (stale reads acceptable)
Failover: Promote secondary to primary

2. Multi-Master Replication:

US-East ↔ EU ↔ APAC
↓ ↓ ↓
Bidirectional replication
All DCs accept writes
Conflict resolution required

3. Consensus-Based (Raft/Paxos):

Quorum writes across DCs
Majority must acknowledge
Strong consistency guarantee
Higher latency for writes

Implementation Considerations:

- Conflict resolution (Last-Write-Wins, CRDTs)
- Cross-DC latency (50–200ms)
- Network partition handling
- Bandwidth optimization (delta sync)

30. SIM Card Store System

Problem Overview

Design a system for a SIM card store to generate unpredictable phone numbers, track sold/unused inventory, and enable efficient retrieval.

Key Requirements

- Generate random phone numbers (no patterns)
- Track status: available, sold, reserved
- Efficient search and allocation
- Prevent number prediction

Schema:

```
phone_numbers (
    id BIGSERIAL PRIMARY KEY,
    phone_number VARCHAR(15) UNIQUE NOT NULL,
    status VARCHAR(20) NOT NULL,
    allocated_at TIMESTAMP,
    customer_id BIGINT,
    random_seed BYTEA -- for generation
);

CREATE INDEX idx_phone_status ON phone_numbers(status)
WHERE status = 'available';
```

Number Generation:

```
import secrets
import hashlib

def generate_phone_number():
    """Generate cryptographically random phone number."""
    # Use secrets for unpredictability
    random_bytes = secrets.token_bytes(8)

    # Hash to get number
    hash_val = int(hashlib.sha256(random_bytes).hexdigest(), 16)

    # Convert to 10-digit phone number
    phone = str(hash_val % 10000000000).zfill(10)

    return f"+1{phone}"
```

31-32. Optimizing Hotel Search Results & Ranking Algorithm

Search Optimization Strategies

1. Inverted Index:

```

hotels_by_location[
    "New York": [hotel1, hotel2, ...],
    "San Francisco": [hotel3, hotel4, ...]
]

hotels_by_amenity[
    "pool": [hotel1, hotel5, ...],
    "wifi": [hotel2, hotel3, ...]
]

```

2. Ranking Algorithm:

```

def calculate_hotel_score(hotel, user_prefs):
    score = 0

    # Price score (inverse - cheaper is better)
    price_score = 1 / (hotel.price / 100 + 1)
    score += price_score * 0.3

    # Rating score
    rating_score = hotel.rating / 5.0
    score += rating_score * 0.4

    # Distance score (closer is better)
    distance_score = 1 / (hotel.distance_km + 1)
    score += distance_score * 0.2

    # Amenity match score
    amenity_score = len(hotel.amenities & user_prefs.amenities) /
    len(user_prefs.amenities)
    score += amenity_score * 0.1

    return score

```

3. Elasticsearch Query:

```
{
  "query": {
    "function_score": {
      "query": {
        "bool": {
          "must": [
            {"geo_distance": {
              "distance": "10km",
              "location": {"lat": 40.7, "lon": -74.0}
            }},
          ]
        }
      }
    }
  }
}
```


```
        {"range": {"price": {"lte": 200}}}
    ]
}
},
"functions": [
    {"field_value_factor": {
        "field": "rating",
        "factor": 2.0
    }},
    {"gauss": {
        "price": {
            "origin": 100,
            "scale": 50
        }
    }}
]
}
}
```

33. Real-time Chat Application

Problem Overview

Design a real-time chat application supporting one-to-one and group messaging with typing indicators, read receipts, and message history.

Architecture



Message Schema:

```
messages (
    conversation_id UUID,
    message_id TIMEUUID,
    sender_id UUID,
    content TEXT,
    timestamp TIMESTAMP,
    status VARCHAR(20), -- sent, delivered, read
    PRIMARY KEY (conversation_id, message_id)
) WITH CLUSTERING ORDER BY (message_id DESC);
```

WebSocket Handler:

```
io.on('connection', (socket) => {
  socket.on('send_message', async (data) => {
    const message = {
      id: uuidv1(),
      conversation_id: data.conversation_id,
      sender_id: socket.user_id,
      content: data.content,
      timestamp: new Date()
    };

    // Store in Cassandra
    await cassandra.insert('messages', message);

    // Publish via Redis Pub/Sub
    redis.publish(`chat:${data.conversation_id}`,
      JSON.stringify(message));

    // Emit to recipients
    io.to(data.conversation_id).emit('new_message', message);
  });
});
```

34. Distributed Message Broker

Problem Overview

Design a high-throughput distributed message broker like Kafka for pub/sub messaging with guaranteed delivery and ordering.

Architecture

Producers → Brokers (Partitioned Topics) → Consumers

Topic: orders
|— Partition 0 → Consumer Group A

- └ Partition 1 → Consumer Group A
- └ Partition 2 → Consumer Group A

Features:

- Partitioning for parallelism
- Replication for durability
- Consumer groups for load balancing
- Offset management for exactly-once

Key Design Decisions:

- Log-based storage (append-only)
- Partition assignment (consistent hashing)
- Replication (ISR - In-Sync Replicas)
- Consumer offset commits

35. Cloud File Storage (Dropbox)

Problem Overview

Design a cloud file storage service with sync, sharing, and version control (covered in detail in Problem #9).

Key Features:

- Chunking (4MB blocks)
- Deduplication (block-level)
- Delta sync (only changed chunks)
- Conflict resolution (version vectors)

36. Distributed Configuration Store

Problem Overview

Design a distributed configuration store like etcd or ZooKeeper for storing application config with strong consistency.

Features

- Key-value storage
- Watch mechanism (notifications on changes)
- TTL support
- Transactions
- Consensus (Raft algorithm)

Use Cases:

- Service discovery
- Leader election
- Distributed locks

- Feature flags

```
# Example usage
config_store = DistributedConfig()

# Set configuration
await config_store.set("/app/db_host", "db.example.com")

# Watch for changes
async for event in config_store.watch("/app/"):
    print(f"Config changed: {event.key} = {event.value}")
```

37. Nearby Places Recommender (Yelp)

Problem Overview

Design a location-based service to find nearby restaurants/businesses with ratings and reviews (similar to Problem #6 - Proximity Search).

Geospatial Indexing:

- Quadtree partitioning
- Geohash (precision-based cells)
- PostGIS (R-tree index)
- Elasticsearch geo queries

Ranking:

```
def rank_places(places, user_location, user_prefs):
    for place in places:
        score = 0

        # Distance penalty
        distance = haversine(user_location, place.location)
        distance_score = 1 / (distance + 0.1)
        score += distance_score * 0.4

        # Rating
        score += (place.rating / 5.0) * 0.3

        # Review count (popularity)
        review_score = min(place.review_count / 1000, 1.0)
        score += review_score * 0.2

        # Price match
        if place.price_range == user_prefs.price_range:
            score += 0.1
```

```

    place.score = score

    return sorted(places, key=lambda p: p.score, reverse=True)

```

38. Gaming Leaderboard

Problem Overview

Design a real-time gaming leaderboard supporting millions of players with efficient rank queries and updates.

Architecture

Player Score Update:

```
ZADD leaderboard:global player_id score
```

Get Player Rank:

```
ZREVRANK leaderboard:global player_id
```

Get Top 100:

```
ZREVRANGE leaderboard:global 0 99 WITHSCORES
```

Get Nearby Players (player's rank ±10):

```
rank = ZREVRANK leaderboard:global player_id
```

```
ZREVRANGE leaderboard:global (rank-10) (rank+10) WITHSCORES
```

Redis Sorted Set Operations:

```

class Leaderboard:
    def __init__(self, redis_client):
        self.redis = redis_client
        self.key = "leaderboard:global"

    def update_score(self, player_id, score):
        """Update player score."""
        self.redis.zadd(self.key, {player_id: score})

    def get_rank(self, player_id):
        """Get player's rank (1-indexed)."""
        rank = self.redis.zrevrank(self.key, player_id)
        return rank + 1 if rank is not None else None

    def get_top_players(self, limit=100):
        """Get top N players."""
        return self.redis.zrevrange(
            self.key, 0, limit - 1, withscores=True
        )

```

```

def get_player_context(self, player_id, context=10):
    """Get nearby players."""
    rank = self.redis.zrevrank(self.key, player_id)
    if rank is None:
        return []

    start = max(0, rank - context)
    end = rank + context

    return self.redis.zrevrange(
        self.key, start, end, withscores=True
    )

```

Sharding for Millions of Players:

Shard by score range:

- Shard 1: scores 0–1000
- Shard 2: scores 1001–2000
- Shard 3: scores 2001–3000
- ...

Or by player_id hash:

```
shard = hash(player_id) % num_shards
```

39. Hotel Reservation System (Detailed)

Covered comprehensively in Problem #15 with double-booking prevention strategies.

40. Multilingual Database Schema

Problem Overview

Design database schema to support multiple languages for product catalogs, content, etc.

Approaches

1. Separate Translation Tables:

```

products (
    id BIGINT PRIMARY KEY,
    sku VARCHAR(50),
    price DECIMAL(10,2)
);

product_translations (
    product_id BIGINT REFERENCES products(id),

```

```

language_code VARCHAR(5),
name VARCHAR(255),
description TEXT,
PRIMARY KEY (product_id, language_code)
);

-- Query
SELECT p.* , pt.name, pt.description
FROM products p
JOIN product_translations pt ON p.id = pt.product_id
WHERE pt.language_code = 'es';

```

2. JSONB Columns:

```

products (
    id BIGINT PRIMARY KEY,
    sku VARCHAR(50),
    price DECIMAL(10,2),
    names JSONB, -- {"en": "Laptop", "es": "Portátil", "fr": "Ordinateur
portable"}
    descriptions JSONB
);

-- Query
SELECT id, names->>'es' as name
FROM products;

```

3. Separate Tables per Language (not recommended):

```

products_en (...)
products_es (...)
products_fr (...)
-- Maintenance nightmare

```

Best Practice: Separate translation tables for normalized data, JSONB for simpler use cases.

41. System Improvement Analysis

Problem Overview

Given an existing badly designed system, identify flaws and propose improvements with technical justifications.

Analysis Framework

1. Identify Issues:

- Performance bottlenecks
- Scalability limitations
- Single points of failure
- Security vulnerabilities
- Poor data modeling
- Lack of monitoring

2. Propose Solutions:

- Caching strategies
- Database indexing
- Load balancing
- Replication/sharding
- API rate limiting
- Observability tools

Example Analysis:

Current System Issues:

1. Single PostgreSQL instance → Add read replicas + connection pooling
2. No caching → Add Redis cache layer (80% hit rate target)
3. Synchronous payment processing → Async queue (Kafka)
4. No rate limiting → Implement token bucket algorithm
5. Large table scans → Add composite indexes on query patterns

Expected Improvements:


- Latency: 2000ms → 200ms (p95)
- Throughput: 100 req/s → 1000 req/s
- Availability: 99% → 99.9%

42. Multi-Property Hotel Management

Problem Overview

Design a platform for managing multiple hotel properties with centralized inventory, bookings, and analytics.

Architecture





Data Isolation:

```

-- Shared schema with tenant_id
bookings (
    id UUID PRIMARY KEY,
    property_id BIGINT NOT NULL,
    room_id BIGINT,
    ...
);


-- Row-level security
CREATE POLICY property_isolation ON bookings
FOR ALL
USING (property_id = current_setting('app.current_property_id')::bigint);
  
```

43. Scalable Android System

Problem Overview

Design a robust Android architecture with caching, offline support, and state management.

Architecture Layers



Caching Strategy:

```

class Repository(
    private val api: ApiService,
    private val cache: LocalCache
) {
    suspend fun getData(): Result<Data> {
        // Try cache first
        val cached = cache.get()
        if (cached != null && !cached.isStale()) {
            return Result.Success(cached)
        }

        // Fetch from network
        return try {
            val fresh = api.fetchData()
            cache.save(fresh)
            Result.Success(fresh)
        } catch (e: Exception) {
            // Fallback to stale cache
            cached?.let { Result.Success(it) }
            ?: Result.Error(e)
        }
    }
}

```

44. Flight Inventory with Metered APIs

Covered comprehensively in Problem #25 with cost optimization and rate limiting strategies.

45. Store Inventory Management

Problem Overview

Design an inventory management system for retail stores with real-time stock tracking, reorder alerts, and multi-location support.

Features

- Stock level tracking
- Automatic reorder points
- Transfer between locations
- Audit trail for stock movements
- Low stock alerts

Schema:

```

products (
    id BIGSERIAL PRIMARY KEY,
    sku VARCHAR(50) UNIQUE,

```

```

name VARCHAR(255),
category VARCHAR(100),
reorder_point INT,
reorder_quantity INT
);

locations (
    id BIGSERIAL PRIMARY KEY,
    name VARCHAR(255),
    address TEXT,
    type VARCHAR(20) -- warehouse, store
);

inventory (
    id BIGSERIAL PRIMARY KEY,
    product_id BIGINT REFERENCES products(id),
    location_id BIGINT REFERENCES locations(id),
    quantity INT NOT NULL DEFAULT 0,
    last_updated TIMESTAMP DEFAULT NOW(),
    UNIQUE(product_id, location_id)
);

stock_movements (
    id BIGSERIAL PRIMARY KEY,
    product_id BIGINT,
    location_id BIGINT,
    movement_type VARCHAR(20), -- in, out, transfer, adjustment
    quantity INT,
    reference_id VARCHAR(100), -- order_id, transfer_id, etc.
    created_at TIMESTAMP DEFAULT NOW()
);

-- Trigger for low stock alerts
CREATE OR REPLACE FUNCTION check_reorder_point()
RETURNS TRIGGER AS $$
BEGIN
    IF NEW.quantity <= (
        SELECT reorder_point FROM products WHERE id = NEW.product_id
    ) THEN
        INSERT INTO reorder_alerts (product_id, location_id, current_quantity)
        VALUES (NEW.product_id, NEW.location_id, NEW.quantity);
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER inventory_reorder_check
AFTER UPDATE OF quantity ON inventory
FOR EACH ROW
EXECUTE FUNCTION check_reorder_point();

```

Inventory Operations:

```
class InventoryService:
    def adjust_stock(self, product_id, location_id, quantity,
                    movement_type, reference_id=None):
        """Adjust stock with audit trail."""
        with db.transaction():
            # Update inventory
            db.execute("""
                INSERT INTO inventory (product_id, location_id, quantity)
                VALUES ($1, $2, $3)
                ON CONFLICT (product_id, location_id)
                DO UPDATE SET
                    quantity = inventory.quantity + EXCLUDED.quantity,
                    last_updated = NOW()
            """, product_id, location_id, quantity)

            # Record movement
            db.execute("""
                INSERT INTO stock_movements
                (product_id, location_id, movement_type, quantity,
                reference_id)
                VALUES ($1, $2, $3, $4, $5)
            """, product_id, location_id, movement_type, quantity,
reference_id)

    def transfer_stock(self, product_id, from_location, to_location,
                       quantity):
        """Transfer stock between locations."""
        with db.transaction():
            # Deduct from source
            self.adjust_stock(product_id, from_location, -quantity,
                              'transfer', f'TXF-{uuid4()}')

            # Add to destination
            self.adjust_stock(product_id, to_location, quantity,
                              'transfer', f'TXF-{uuid4()}')

    def get_low_stock_items(self, location_id=None):
        """Get items below reorder point."""
        query = """
            SELECT p.*, i.quantity, i.location_id
            FROM products p
            JOIN inventory i ON p.id = i.product_id
            WHERE i.quantity <= p.reorder_point
        """

        if location_id:
            query += " AND i.location_id = $1"
            return db.fetch(query, location_id)

        return db.fetch(query)
```

Summary of All 45 Solutions

This document provides comprehensive system design solutions covering:

Core Systems (1-15):

- Streaming platforms, search systems, file storage
- Booking systems with concurrency control
- Distributed scheduling, payment gateways

Concepts & Patterns (16-23):

- Caching strategies, sharding, consistency models
- Rate limiting, top-K algorithms

Specialized Systems (24-45):

- Reconciliation, flight inventory
- Distributed KV stores, message brokers
- Real-time chat, leaderboards
- Multi-tenant platforms

Key Themes:

- Horizontal scalability via partitioning/sharding
- High availability through replication
- Low latency via multi-level caching
- Strong consistency where needed (transactions)
- Eventual consistency where acceptable (performance)
- Observability and monitoring throughout

Each solution includes: Architecture diagrams Database schemas Implementation code

Trade-off analysis Scalability considerations

End of Complete System Design Solutions (1-45)